

# Distributed Memory Programming (2D)

## ICS632: Principles of High Performance Computing

Henri Casanova (henric@hawaii.edu)

Fall 2015


# Outline

1 Introduction

2 Matrix Multiplication

3 Conclusion

# Grid/Torus of Processors



## 2-D Data Distribution

- We'll only consider 2-D *square* matrices
- There is thus a natural “block” data distribution:

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	$A_{0,3}$	$A_{0,4}$
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	$A_{1,3}$	$A_{1,4}$
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	$A_{2,3}$	$A_{2,4}$
$A_{3,0}$	$A_{3,1}$	$A_{3,2}$	$A_{3,3}$	$A_{3,4}$
$A_{4,0}$	$A_{4,1}$	$A_{4,2}$	$A_{4,3}$	$A_{4,4}$

- Each of the  $p$  processes holds a  $N/p \times N/p$  matrix block of an  $N \times N$  matrix

# How do Matrices Get Distributed?

- You can do whatever you want, but what about libraries?
- Option #1 - Centralized: when calling a function (e.g., matrix multiplication) the input data is available on a single “master” machine (perhaps in a file) the input data must then be distributed among workers the output data must be undistributed and returned to the “master” machine (perhaps in a file)
  - More natural/easy for the user
  - The library makes data distribution decisions transparently
  - Prohibitively expensive if one does sequences of operations!!
- Option #2 - Distributed: When calling a function (e.g., matrix multiplication), one assumes that the input is already distributed and the output is left distributed
  - More work for the user
  - May lead to “redistribution” of data in between calls, which is harder for the user and may be costly
- Most current software adopt the distributed approach
- We always assume that the data is magically already distributed

# Outline

1 Introduction

**2 Matrix Multiplication**

3 Conclusion

# Outer-Product Algorithm

- Let's see one classic matrix-multiplication algorithm
- Consider the k-i-j order:

## k-i-j Matrix Multiplication



```
for (k=0; k < N; k++)
  for (i=0; i < N; i++)
    for (j=0; j < N; j++)
      C[i][j] += A[i][k] * B[k][j];
```

- This is a sequence of  $N$  outer-products!
  - Multiply a column vector by a row vector

# Matrix Multiplication: $N$ Outer-Products

## k-j-j Matrix Multiplication

```
for (k=0; k < N; k++)
    // Multiply a column of A by a row of B
    for (i=0; i < N; i++)
        for (j=0; j < N; j++)
            C[i][j] += A[i][k] * B[k][j];
```






# Matrix Multiplication: $N$ Outer-Products

## k-j-j Matrix Multiplication

```
for (k=0; k < N; k++)
    // Multiply a column of A by a row of B
    for (i=0; i < N; i++)
        for (j=0; j < N; j++)
            C[i][j] += A[i][k] * B[k][j];
```




# Matrix Multiplication: $N$ Outer-Products

## k-j-j Matrix Multiplication

```

for (k=0; k < N; k++)
    // Multiply a column of A by a row of B
    for (i=0; i < N; i++)
        for (j=0; j < N; j++)
            C[i][j] += A[i][k] * B[k][j];
    
```




# Matrix Multiplication: $N$ Outer-Products

## k-j-j Matrix Multiplication

```

for (k=0; k < N; k++)
    // Multiply a column of A by a row of B
    for (i=0; i < N; i++)
        for (j=0; j < N; j++)
            C[i][j] += A[i][k] * B[k][j];
    
```




# Matrix Multiplication: $N$ Outer-Products

## k-j-j Matrix Multiplication

```

for (k=0; k < N; k++)
    // Multiply a column of A by a row of B
    for (i=0; i < N; i++)
        for (j=0; j < N; j++)
            C[i][j] += A[i][k] * B[k][j];
    
```




# Matrix Multiplication: $N$ Outer-Products

## k-j-j Matrix Multiplication

```

for (k=0; k < N; k++)
    // Multiply a column of A by a row of B
    for (i=0; i < N; i++)
        for (j=0; j < N; j++)
            C[i][j] += A[i][k] * B[k][j];
    
```



# So What??

- Why do we care about thinking of matrix multiplication this way???
- Note that in principles there are  $n^3!$  possible sequential algorithms
- Because it's possible to have a very elegant parallel algorithm
- Let's see a small example for a  $4 \times 4$  grid of processors...

# The Outer-Product Algorithm

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	$A_{0,3}$	$A_{0,4}$
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	$A_{1,3}$	$A_{1,4}$
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	$A_{2,3}$	$A_{2,4}$
$A_{3,0}$	$A_{3,1}$	$A_{3,2}$	$A_{3,3}$	$A_{3,4}$
$A_{4,0}$	$A_{4,1}$	$A_{4,2}$	$A_{4,3}$	$A_{4,4}$

$B_{0,0}$	$B_{0,1}$	$B_{0,2}$	$B_{0,3}$	$B_{0,4}$
$B_{1,0}$	$B_{1,1}$	$B_{1,2}$	$B_{1,3}$	$B_{1,4}$
$B_{2,0}$	$B_{2,1}$	$B_{2,2}$	$B_{2,3}$	$B_{2,4}$
$B_{3,0}$	$B_{3,1}$	$B_{3,2}$	$B_{3,3}$	$B_{3,4}$
$B_{4,0}$	$B_{4,1}$	$B_{4,2}$	$B_{4,3}$	$B_{4,4}$

$C_{0,0}$	$C_{0,1}$	$C_{0,2}$	$C_{0,3}$	$C_{0,4}$
$C_{1,0}$	$C_{1,1}$	$C_{1,2}$	$C_{1,3}$	$C_{1,4}$
$C_{2,0}$	$C_{2,1}$	$C_{2,2}$	$C_{2,3}$	$C_{2,4}$
$C_{3,0}$	$C_{3,1}$	$C_{3,2}$	$C_{3,3}$	$C_{3,4}$
$C_{4,0}$	$C_{4,1}$	$C_{4,2}$	$C_{4,3}$	$C_{4,4}$

```

for (k=0; k < N; k++)
  for (i=0; i < N; i++)
    for (j=0; j < N; j++)
      // Block operations
      C_{i,j} += A_{i,k} * B_{k,j};

```



- At step  $k$ , processor  $(i, j)$  needs  $A_{i,k}$  and  $B_{k,j}$
- If  $j = k$ , processor  $(i, j)$  already has  $A_{i,k}$ ,  
otherwise it must receive it from processor  $(i, k)$
- If  $i = k$ , processor  $(i, j)$  already has  $B_{k,j}$ ,  
otherwise it must receive it from processor  $(k, j)$

# Communication Pattern

- From the previous slide:
  - At step  $k$ , processor  $(i, j)$  needs  $A_{i,k}$  and  $B_{k,j}$
  - If  $j = k$ , processor  $(i, j)$  already has  $A_{i,k}$ , otherwise it must receive it from processor  $(i, k)$
  - If  $i = k$ , processor  $(i, j)$  already has  $B_{k,j}$ , otherwise it must receive it from processor  $(k, j)$
- Therefore, at step  $k = 0, \dots, p - 1$ :
  - $\forall i$ , processor  $(i, k)$  broadcasts its block of  $A$  to all processors in row  $i$
  - $\forall j$ , processor  $(k, j)$  broadcasts its block of  $B$  to all processors in column  $j$
- Let's see it on a picture...



# Communication Pattern: $k = 1$



# The Outer-Product Algorithm

```
p = sqrt(num_procs());
int A[N/p][N/p], B[N/p][N/p], C[N/p][N/p];
int bufferA[N/p][N/p], bufferB[N/p][N/p];

(myrow, mycol) = my_2D_rank();

for (int k=0; k < p; k++) {
    // Broadcast A along rows
    BroadcastRow(myrow, k), A, bufferA, N/p * N/p);
    // Broadcast B along columns
    BroadcastColumn(k, mycol), B, bufferB, N/p * N/p);

    // Multiply Matrix blocks (assuming a convenient MatrixMultiplyAdd() function)
    if ((myrow == k) && (mycol == k))
        MatrixMultiplyAdd(C, A, B, N/p, N/p);
    else if (myrow == k)
        MatrixMultiplyAdd(C, bufferA, B, N/p, N/p);
    else if (mycol == k)
        MatrixMultiplyAdd(C, A, bufferB, N/p, N/p);
    else
        MatrixMultiplyAdd(C, bufferA, bufferB, N/p, N/p);
}
```

# Performance Analysis

- $\beta$ : time to do a row/column broadcast
- $\gamma$ : time to compute a block
- No overlap: time =  $p \times (\beta + \beta + \gamma)$
- Some overlap: time =  $\beta + \beta + (p - 1) \max(\beta + \beta, \gamma) + \gamma$
- This algorithm is in fact asymptotically optimal

# Grid vs. Ring

- On a Ring, with a 1-D data distribution we already have an asymptotically optimal algorithm (same ideas as for the 1-D Matrix Vector multiply)
- So who cares about this more complex algorithm?
- If  $N$  is huge, we don't care
- But in fact, using a 2-D distribution reduces communication costs
- The algorithm sends less data overall
- And it can be proven that even if the underlying platform is not a torus, the 2-D algorithm is still better than the 1-D algorithm


# Other Matrix Multiplication Algorithms

- People have come up with many algorithms
  - Cannon (1969)
  - Fox (1987)
  - Snyder (1992)
  - ...
- They all correspond to “cruising” through the operations in different (possibly really confusing) order
- Some begin by shuffling things around in each matrix!


## 2-D Block Cyclic Distribution

- If  $N \gg p$ , then blocks are large
- This can be a problem as it limits parallelism
- One “swiss army knife” solution is to use a **2-D Block Cyclic Distribution**
  - Doesn't matter which way the computation “moves”, we should be ok...
- Let's see what that looks like...


## 2-D Block Cyclic Distribution



## 2-D Block Cyclic Distribution




$2 \times 3$  processor grid







## 2-D Block Cyclic Distribution




$2 \times 3$  processor grid




## 2-D Block Cyclic Distribution




$2 \times 3$  processor grid




## 2-D Block Cyclic Distribution




$2 \times 3$  processor grid




## 2-D Block Cyclic Distribution



$2 \times 3$  processor grid





# What One Processor Holds



Fun local/global indices computations!

Processor (0,1)



```
int A[N/2][N/3];
```

## 2-D Block Cyclic Distribution

- All algorithms we've seen so far can be implemented with the 2-D Block Cyclic distribution
- If you don't abstract it away a bit, the code becomes horrendous
- But you'd get performance benefits from it
- For instance, the ScaLAPACK library recommends 2-D block cyclic distributions
  - It actually supports other distributions

# Outline

1 Introduction

2 Matrix Multiplication

**3 Conclusion**

# What was all this????

- These lecture notes are representative of “traditional” parallel computing
- Similar algorithms have been studied for decades, their performance analyzed in depth
  - Using various models/assumptions
- We have a programming assignment on this topic...
- The main caveat of all this material is that when the assumptions break down, or when the platform becomes is very complex, then many difficulties arise
- For instance, when the platform is heterogeneous...