

# Fault Tolerance (and a bit of Energy)

## ICS632: Principles of High-Performance Computing

Henri Casanova (henric@hawaii.edu)

Fall 2015

# Disclaimer

- These slides will provide a quick overview of many important questions faced by theoreticians and practitioners in HPC
- Most of this is still very much in the research and development stage
- It would take us way too long to explore all this content in depth
- We'll learn about some of the topics through research paper presentations

# Towards exascale platforms

- The scale of parallel platforms has been steadily increasing for decades
  - number of cabinets, number of blades, number of processors, number of cores
- Several challenges arise (topics of entire conferences/workshops!)
  - Resilience to failures
  - Power consumption
  - Heat management
  - Network interconnects
- Some for people who build platforms, some for people who use platforms

# Towards exascale platforms

- The scale of parallel platforms has been steadily increasing for decades
  - number of cabinets, number of blades, number of processors, number of cores
- Several challenges arise (topics of entire conferences/workshops!)
  - Resilience to failures
  - Power consumption
  - Heat management
  - Network interconnects
- Some for people who build platforms, **some for people who use platforms**

# Outline

1 Why Fault Tolerance?

2 Checkpointing

3 A little theory

4 A little practice

5 Energy Consumption

# Faults, Errors, and Failures

- A **fault** corresponds to an execution that does not go according to specification (hardware, software)
- A fault may cause an **error**, i.e., an incorrect system state
- An error may cause a **failure**, i.e., an unacceptable behavior (e.g., crash, wrong output)
- WARNING: terminology is all over the place in the literature
- Why do failures occur? Because redundancy is too costly
- General question: how to achieve reliability out of unreliable components?

# All kinds of faults/failures

- Hardware faults:
  - Detected and corrected by hardware (ECC works)
  - Detected in hardware, but flagged as not corrected (e.g., double bit-flip)
  - Silent (bits are wrong, but you don't know!)
- Software faults:
  - Pure software faults (bugs)
  - Mis-handling of hardware faults
  - OS/firmware faults
  - Many of these are not silent, but not all
- Some failures can be handled by a "reboot", others by a "replace hardware"
  - In practice, "replacing" means "boot a spare"
  - So the two are very similar

## How often do failures occur?

- Say you buy a component with Mean Time Between Failures (MTBF)  $T$
  - Now you put  $n$  components together in a platform
  - The platform's MTBF is  $T/n$
  - When  $n$  is large, the MTBF will be low
  - So we have a problem for (exa)scale
- 
- Failure distributions are often assumed i.i.d. and exponential (memoryless) in the literature
  - But in practice they are likely not i.i.d. and not memoryless (e.g., Weibull), which is much more difficult to work with for theoreticians



# Outline

- 1 Why Fault Tolerance?
- 2 Checkpointing**
- 3 A little theory
- 4 A little practice
- 5 Energy Consumption


# Checkpoint time, recovery time, down time

- Checkpoint time:  $C$  seconds
  - Depends on checkpoint size, storage medium
- Recovery time:  $R$  seconds
  - Depends on checkpoint size, storage medium
- Downtime:  $D$  seconds
  - After a failure, time for the *logical* host to start a recovery
  - Typically: constant time to bring a spare host on-line
- On a single processor:




# Checkpoint time, recovery time, down time

- Checkpoint time:  $C$  seconds
  - Depends on checkpoint size, storage medium
- Recovery time:  $R$  seconds
  - Depends on checkpoint size, storage medium
- Downtime:  $D$  seconds
  - After a failure, time for the *logical* host to start a recovery
  - Typically: constant time to bring a spare host on-line
- On a single processor:



# Coordinated checkpointing on multiple processors



# Outline

- 1 Why Fault Tolerance?
- 2 Checkpointing
- 3 A little theory**
- 4 A little practice
- 5 Energy Consumption

# “Theoretical” questions


- How often should one checkpoint?
  - Goal: minimize expected makespan
- Should checkpointing be periodic?
  - What everybody has done in practice
- Let's look at some development on one processor to get a sense of what theoreticians have done

## Assuming Periodic checkpointing

- Many authors have proposed analytical expressions for the “optimal” checkpointing frequency, **assuming that checkpointing must be periodic**
  - or approximations thereof
- A famous such expression is given by J.W. Young
  - *A First Order Approximation to the Optimal Checkpoint Interval*, Communication of the ACM, 1974.
- Let's review how this bound is achieved, which will provide context for recent results

# Young's approach

- Let  $T$  be the time between checkpoints
- Let  $T_w$  be the compute time wasted due to a failure
- Let us consider an interval  $T_i$  in between two failures



$$T_i = n \times (T + C) + T_w$$



# Young's assumptions

- Objective: minimize  $\mathbb{E}[T_w]$ , where  $T_w = T_i - n(T + C)$
- Four questionable assumptions:
  - Assumption #1: failure arrival times follow an Exponential distribution of mean  $1/\lambda$
  - Assumption #2: failures do not occur during checkpointing
  - Assumption #3: failures do not occur during recovery
  - Assumption #4: failures do not occur during downtime (i.e., on another processor)

# Obtaining the approximation

- If  $T_i$  is in between  $n(T + C)$  and  $(n + 1)(T + C)$ , then  
 $T_w = T_i - nT$ 
  - subtract from the whole time the “useful” compute time
- Therefore:

$$\mathbb{E}[T_w] = \sum_{n=0}^{\infty} \int_{n(T+C)}^{(n+1)(T+C)} (t - nT) \lambda e^{\lambda t} dt$$

$$\Rightarrow \quad \dots \quad \Rightarrow \quad \mathbb{E}[T_w] = 1/\lambda + T/(1 - e^{\lambda(T+C)})$$

$$\Rightarrow \quad \frac{d\mathbb{E}[T_w]}{dT} = \frac{1 - e^{\lambda(T+C)} + T e^{\lambda(T+C)}}{(1 - e^{\lambda(T+C)})^2}$$

## Obtaining the approximation

- To find the optimal  $T$  we solve:  $e^{\lambda T} e^{\lambda C} (1 - \lambda T) - 1 = 0$
- $\lambda T$  is small (because  $1/\lambda$  is large)
- Using the approx  $e^{\lambda T} = 1 + \lambda T + \lambda T^2/2$  and neglecting terms of degree 3, we obtain  $\frac{1}{2}(\lambda T)^2 = 1 - e^{-\lambda C}$
- $\lambda C$  is even smaller (because  $C$  is small), and we use the approx  $e^{-\lambda C} = 1 - \lambda C$
- We obtain Young's approximation:

$$T_{opt} \sim \sqrt{2C/\lambda}$$

- Example:  $MTBF = 48$  hours,  $C = 1$  minute, then checkpoint every  $T_{opt} = 75.9$  minutes

## Going further

- With recovery time, Young's approximation becomes

$$T_{opt} \sim \sqrt{2C(R + 1/\lambda)}$$

- J.T. Daly, *A higher order estimate of the optimum checkpoint interval for restart dumps*, FGCS 2006,

- Daly goes further and proposes the following approximation:

$$T_{opt} = \sqrt{2C/\lambda} \left[ 1 + \sqrt{\frac{C\lambda}{18}} + \frac{C\lambda}{18} \right] - C, \quad \text{if } C < 2/\lambda$$

$$T_{opt} = 1/\lambda, \quad \text{if } C \geq 2/\lambda$$

- Example:  $MTBF = 48$  hours,  $C = 1$  minute, then checkpoint every  $T_{opt} = 75.2$  minutes

# Daly's bound

- Pretty close to Young's bound unless  $C$  is relatively large
- Daly doesn't ignore recovery (even though in the final formula  $R$  isn't there!)
- Daly estimates the fraction of wasted work once a failure occurs ( $T_w$ ) much better
- Daly allows failures during recovery, **but not during checkpointing**

# Research questions

- The assumption that  $C \ll 1/\lambda = MTBF$  is likely invalid at large scale
  - $MTBF_{platform} = MTBF_{proc}/\#procs$
- The assumption that  $C$  is so small that there are no failures during checkpointing is likely invalid at large scale
- The assumption that there cannot be a failure during a downtime (of another host) is likely invalid at large scale
  - “Cascading” failures
- Are we even sure that periodic is optimal???

# Research questions

- The assumption that  $C \ll 1/\lambda = MTBF$  is likely invalid at large scale
  - $MTBF_{platform} = MTBF_{proc}/\#procs$
- The assumption that  $C$  is so small that there are no failures during checkpointing is likely invalid at large scale
- The assumption that there cannot be a failure during a downtime (of another host) is likely invalid at large scale
  - “Cascading” failures
- Are we even sure that periodic is optimal???

# Research questions

- The assumption that  $C \ll 1/\lambda = MTBF$  is likely invalid at large scale
  - $MTBF_{platform} = MTBF_{proc}/\#procs$
- The assumption that  $C$  is so small that there are no failures during checkpointing is likely invalid at large scale
- The assumption that there cannot be a failure during a downtime (of another host) is likely invalid at large scale
  - “Cascading” failures
- Are we even sure that periodic is optimal???



# Research questions

- The assumption that  $C \ll 1/\lambda = MTBF$  is likely invalid at large scale
  - $MTBF_{platform} = MTBF_{proc} / \#procs$
- The assumption that  $C$  is so small that there are no failures during checkpointing is likely invalid at large scale
- The assumption that there cannot be a failure during a downtime (of another host) is likely invalid at large scale
  - “Cascading” failures
- Are we even sure that periodic is optimal???

# New Problem statement and definition

- We focus on  $C_{max}$  minimization, or rather  $\mathbb{E}[C_{max}]$  minimization
  - After all, this is the real objective, not  $\mathbb{E}[T_w]$
- Let us first study the case of a *sequential* job that starts at time  $t_0$ 
  - Sounds simple, but in fact it's already quite complicated
- Let's not make any assumption on the distribution for now: starting at time  $t_0$ , failures occur at time  $t_n = t_0 + \sum_{m=1}^n X_m$ , where the  $X_m$ 's are *i.i.d* random variable
- $P_{suc}(x|\tau)$ : the probability that there is no failure for the next  $x$  seconds knowing that the last failure was  $\tau$  seconds ago
  - Given by the probability distribution



# New Problem statement and definition

- We focus on  $C_{max}$  minimization, or rather  $\mathbb{E}[C_{max}]$  minimization
  - After all, this is the real objective, not  $\mathbb{E}[T_w]$
- Let us first study the case of a *sequential* job that starts at time  $t_0$ 
  - Sounds simple, but in fact it's already quite complicated
- Let's not make any assumption on the distribution for now: starting at time  $t_0$ , failures occur at time  $t_n = t_0 + \sum_{m=1}^n X_m$ , where the  $X_m$ 's are *i.i.d* random variable
- $P_{suc}(x|\tau)$ : the probability that there is no failure for the next  $x$  seconds knowing that the last failure was  $\tau$  seconds ago
  - Given by the probability distribution

# New Problem statement and definition

- We focus on  $C_{max}$  minimization, or rather  $\mathbb{E}[C_{max}]$  minimization
  - After all, this is the real objective, not  $\mathbb{E}[T_w]$
- Let us first study the case of a *sequential* job that starts at time  $t_0$ 
  - Sounds simple, but in fact it's already quite complicated
- Let's not make any assumption on the distribution for now: starting at time  $t_0$ , failures occur at time  $t_n = t_0 + \sum_{m=1}^n X_m$ , where the  $X_m$ 's are *i.i.d* random variable
- $P_{suc}(x|\tau)$ : the probability that there is no failure for the next  $x$  seconds knowing that the last failure was  $\tau$  seconds ago
  - Given by the probability distribution

# Problem statement and definition

- Let  $w$  denote an *amount of work* that remains to be done
  - i.e., a number of seconds of computation
- Let  $T(w|\tau)$  be the time needed to complete  $w$  units of work given that the last failure was  $\tau$  seconds ago
  - Accounting for failures
- Our objective: minimize  $\mathbb{E}[T(W_{total}|\tau)]$ 
  - $W_{total}$ : the total amount of work to be done
  - $\tau$ : the number of seconds since the last failure at  $t_0$

# Problem statement and definition

- Let  $w$  denote an *amount of work* that remains to be done
  - i.e., a number of seconds of computation
- Let  $T(w|\tau)$  be the time needed to complete  $w$  units of work given that the last failure was  $\tau$  seconds ago
  - Accounting for failures
- Our objective: minimize  $\mathbb{E}[T(W_{total}|\tau)]$ 
  - $W_{total}$ : the total amount of work to be done
  - $\tau$ : the number of seconds since the last failure at  $t_0$

# Problem statement and definition

- Let  $w$  denote an *amount of work* that remains to be done
  - i.e., a number of seconds of computation
- Let  $T(w|\tau)$  be the time needed to complete  $w$  units of work given that the last failure was  $\tau$  seconds ago
  - Accounting for failures
- Our objective: minimize  $\mathbb{E}[T(W_{total}|\tau)]$ 
  - $W_{total}$ : the total amount of work to be done
  - $\tau$ : the number of seconds since the last failure at  $t_0$



# Checkpointing strategy

- A **checkpointing strategy** is a decision procedure as follows
- Given  $w$  and  $\tau$ , how much work  $w_1$  should we attempt?
  - The attempted amount of work is called a “chunk”
- *Attempt*: repeatedly try the chunk until success
  - Success:  $w_1 + C$  seconds without failure (note the  $+C$ )
- Then, we ask the question again for  $w - w_1$  work and an updated  $\tau$ , until remains 0 units of work
- The checkpointing strategy chooses a sequence of chunk sizes and the number of chunks

# Recursion for $T(w|\tau)$

- $T(0|\tau) = 0$  (no work is done in 0 seconds)
- $T(w|\tau) = w_1 + C + T(w - w_1|\tau + w_1 + C)$ , if there is no failure in the next  $w_1 + C$  seconds
  - Everything went well, we now have  $w - w_1$  work to do, and the last failure is now  $w_1 + C$  seconds further in the past
- $T(w|\tau) = T_{wasted}(w_1 + C|\tau) + T(w|R)$ , otherwise
  - We've wasted a bunch of time, we still have  $w$  work to do, and the last failure happened (ended) right before the last successful recovery
  - $T_{wasted}(w_1 + C|\tau)$ : computation up to a failure + downtime + a recovery during which there can be failures
- We can weigh each case in the recursion by its probability

...

# Computing $\mathbb{E}[T]$

- Probability that there is no failure in the next  $w_1 + C$  second:  $P_{suc}(w_1 + C|\tau)$

- Therefore:

$$\begin{aligned} \mathbb{E}[T(W_{total}|\tau)] = & P_{suc}(w_1 + C|\tau) \times (w_1 + C + \mathbb{E}[T(W_{total} - w_1|\tau + w_1 + C)]) + \\ & (1 - P_{suc}(w_1 + C|\tau)) \times (\mathbb{E}[T_{wasted}(w_1 + C|\tau)] + E(T(W_{total}|R))) \end{aligned}$$

- Remains to compute  $T_{wasted}$

# Computing $\mathbb{E}[T]$

- Probability that there is no failure in the next  $w_1 + C$  second:  $P_{suc}(w_1 + C|\tau)$


- Therefore:

$$\begin{aligned} \mathbb{E}[T(W_{total}|\tau)] = & P_{suc}(w_1 + C|\tau) \times (w_1 + C + \mathbb{E}[T(W_{total} - w_1|\tau + w_1 + C)]) + \\ & (1 - P_{suc}(w_1 + C|\tau)) \times (\mathbb{E}[T_{wasted}(w_1 + C|\tau)] + E(T(W_{total}|R))) \end{aligned}$$

- Remains to compute  $T_{wasted}$

# Computing $T_{wasted}$ , sort of

- $T_{wasted}(w_1 + C|\tau) = T_{lost}(w_1 + C|\tau) + T_{rec}$ 
  - $T_{lost}(x|\tau)$ : amount of time before a failure knowing that a failure occurs in the next  $x$  seconds and that the last failure was  $\tau$  seconds ago
  - $T_{rec}$ : time spent to do the recovery ( $D + R$  in the best case, possibly more if failures during recovery)



# Computing $T_{rec}$

- We can compute  $T_{rec}$  as a function of  $T_{lost}$ :

$$T_{rec} = \begin{cases} D + R & \text{with probability } P_{suc}(R|0), \\ D + T_{lost}(R|0) + T_{rec} & \text{with probability } 1 - P_{suc}(R, 0). \end{cases}$$

- If there is no failure for  $R$  seconds right after the downtime (probability  $P_{suc}(R|0)$ ), then the recovery takes time  $D + R$
- If there is a failure (probability  $1 - P_{suc}(R, 0)$ ), then we spend  $D$  seconds of downtime, waste  $T_{lost}(R|0)$  seconds trying a recovery that would have lasted  $R$  seconds if successful, and then we still have to recover anyway, which requires  $T_{rec}$  seconds

# Computing $T_{rec}$

- Weighing both cases by their probabilities we have

$$\mathbb{E}[T_{rec}] = P_{suc}(R|0) \times (D + R) + \\ (1 - P_{suc}(R|0)) \times (D + \mathbb{E}[T_{lost}(R|0)] + \mathbb{E}[T_{rec}])$$

which gives us:

$$\mathbb{E}[T_{rec}] = D + R + \frac{1 - P_{suc}(R|0)}{P_{suc}(R|0)} (D + \mathbb{E}(T_{lost}(R|0)))$$

and thus

$$\mathbb{E}[T_{wasted}(w_1 + C|\tau)] = \\ \mathbb{E}[T_{lost}(w_1 + C|\tau)] + D + R + \frac{1 - P_{suc}(R|0)}{P_{suc}(R|0)} (D + \mathbb{E}(T_{lost}(R|0)))$$

## Putting it all together, $\mathbb{E}[T(W_{total}, \tau)]$

$$\begin{aligned} \mathbb{E}[T(W_{total}|\tau)] = & P_{suc}(w_1 + C|\tau) \times (w_1 + C + \mathbb{E}[T(W_{total} - w_1|\tau + w_1 + C)]) + \\ & (1 - P_{suc}(w_1 + C|\tau)) \times (\mathbb{E}[T_{lost}(w_1 + C|\tau)] + D + R + \\ & \frac{1 - P_{suc}(R|0)}{P_{suc}(R|0)} (D + \mathbb{E}[T_{lost}(R|0)])) + \mathbb{E}[T(W_{total}|R)] \end{aligned}$$

Easy, right? ☺

- Goal: find  $w_1$  that minimizes the above expression
  - Note the recursion (yikes!)
- Remains to know how to compute  $P_{suc}(x|y)$  and  $\mathbb{E}[T_{lost}(x|y)]$
- And so we make assumptions about the failure distribution...



# Exponential failures

- With exponentially distributed inter-failure times ( $P(X = t) = \lambda e^{-\lambda t}$ ) we obtain:

$$\mathbb{E}[T_{lost}(x|y)] = \int_0^\infty tP(X = t|t < x)dt = \frac{1}{\lambda} - \frac{x}{e^{\lambda x} - 1}$$

$$P_{suc}(x|y) = e^{-\lambda x}$$

- Both expressions above do not involve  $y$  because the Exponential distribution is memoryless
- So we can remove all the “ $\tau$ ” in all probabilities or expectations to simplify notations

# Exponential failures

- With exponentially distributed inter-failure times ( $P(X = t) = \lambda e^{-\lambda t}$ ) we obtain:

$$\mathbb{E}[T_{lost}(x|y)] = \int_0^\infty tP(X = t|t < x)dt = \frac{1}{\lambda} - \frac{x}{e^{\lambda x} - 1}$$

$$P_{suc}(x|y) = e^{-\lambda x}$$

- Both expressions above do not involve  $y$  because the Exponential distribution is memoryless
- So we can remove all the “*tau*” in all probabilities or expectations to simplify notations

# Exponential failures

$$\begin{aligned} \mathbb{E}[T(W_{total})] = & e^{-\lambda(w_1+C)}(w_1 + C + \mathbb{E}[T(W_{total} - w_1)]) + (1 - e^{-\lambda(w_1+C)})(\frac{1}{\lambda} - \\ & \frac{w_1+C}{e^{\lambda(w_1+C)}-1} + D + R + \frac{1-e^{-\lambda R}}{e^{-\lambda R}}(D + \frac{1}{\lambda} - \frac{R}{e^{\lambda R}-1} + \mathbb{E}[T(W_{total})])) \end{aligned}$$

- Assume that there are  $K$  chunks
- We can write an equation for  $\mathbb{E}[T(W_{total})]$  as a function of  $\mathbb{E}[T(W_{total} - w_1)]$
- We can write an equation for  $\mathbb{E}[T(W_{total} - w_1)]$  as a function of  $\mathbb{E}[T(W_{total} - w_1 - w_2)]$
- ...
- We can then solve the recursion!!
  - A LOT of (easy but very tedious) math

# Solved recursion

- We obtain a general form for  $\mathbb{E}[T(W_{total})]$ :

$$\mathbb{E}[T(W_{total})] = A \times \sum_{i=1}^K (e^{\lambda(w_i+C)} - 1)$$

- $e^{\lambda(w_i+C)}$  is a convex function of  $w_i$
- Therefore,  $\mathbb{E}[T(W_{total})]$  is minimized when all  $w_i$ 's are equal
- After decades of periodic checkpointing research, we finally know that it's optimal!! (for exponential failures)
- Important: we made NO approximations
- *Checkpointing Strategies for Parallel Jobs*, Bougeret, Casanova, Rabie, Robert, Vivien [Supercomputing 2011]

# Solved recursion

- We obtain a general form for  $\mathbb{E}[T(W_{total})]$ :

$$\mathbb{E}[T(W_{total})] = A \times \sum_{i=1}^K (e^{\lambda(w_i+C)} - 1)$$

- $e^{\lambda(w_i+C)}$  is a convex function of  $w_i$
- Therefore,  $\mathbb{E}[T(W_{total})]$  is minimized when all  $w_i$ 's are equal
- After decades of periodic checkpointing research, we finally know that it's optimal!! (for exponential failures)
- Important: we made NO approximations
- *Checkpointing Strategies for Parallel Jobs*, Bougeret, Casanova, Rabie, Robert, Vivien [Supercomputing 2011]

# Solved recursion

- We obtain a general form for  $\mathbb{E}[T(W_{total})]$ :

$$\mathbb{E}[T(W_{total})] = A \times \sum_{i=1}^K (e^{\lambda(w_i+C)} - 1)$$

- $e^{\lambda(w_i+C)}$  is a convex function of  $w_i$
- Therefore,  $\mathbb{E}[T(W_{total})]$  is minimized when all  $w_i$ 's are equal
- After decades of periodic checkpointing research, we finally know that it's optimal!! (for exponential failures)
- Important: we made NO approximations
- *Checkpointing Strategies for Parallel Jobs*, Bougeret, Casanova, Rabie, Robert, Vivien [Supercomputing 2011]

## Other theoretical results, and impact

- The same result holds for parallel jobs
  - But we cannot compute the optimal makespan
- If failures are non-memoryless, then periodic is not necessarily optimal
  - The optimal checkpointing strategy can be computed via complicated dynamic programming for Weibull failures
- Nice theory, but if you're a practitioner:
  - You are doing periodic anyway because it's easy
    - Empirical results show that periodic is really bad only in corner cases
  - You know that failures are not i.i.d. anyway, so the theory doesn't apply
  - And if it did, it would be too complex (and slow to compute!)

## Other theoretical results, and impact

- The same result holds for parallel jobs
  - But we cannot compute the optimal makespan
- If failures are non-memoryless, then periodic is not necessarily optimal
  - The optimal checkpointing strategy can be computed via complicated dynamic programming for Weibull failures
- Nice theory, but if you're a practitioner:
  - You are doing periodic anyway because it's easy
    - Empirical results show that periodic is really bad only in corner cases
  - You know that failures are not i.i.d. anyway, so the theory doesn't apply
  - And if it did, it would be too complex (and slow to compute!)



# Outline

- 1 Why Fault Tolerance?
- 2 Checkpointing
- 3 A little theory
- 4 A little practice**
- 5 Energy Consumption

# Checkpointing overhead

- Assume you have some magical, fast way to determine the optimal checkpointing policy for realistic failures
- That's all well and good, but you cannot scale due to  $C$  being large!
  - Parallel efficiency drops as  $p$  increases (Amdahl's law)
- At large scale, your application will spend more time saving state than computing state!!!
- Crucial research problem: reduce checkpointing overhead
- Let's see what people have done/proposed

# Reducing the checkpoint size

- Brute-force **system-level checkpointing**: save the whole address space
  - What Sean is trying to install on our Cray cluster for sequential jobs
- Problem: only a (small) fraction of the address space is needed for a recovery
  - Example: only the iteration number and 2 arrays
- More scalable **application-level checkpointing**: save only what data is necessary for recovery
- Drawback: you must modify application code to checkpoint
  - Note that automatic system-level checkpointing of MPI applications is tough anyway

## Zero checkpoint size?

- In some (lucky?) cases, the algorithm in the application can provide clever ways to “reconstruct” a checkpoint without ever saving it!
- **Algorithm Based Fault Tolerance (ABFT):**
  - A technique by which the algorithm is modified, and made less efficient, to compute on encoded (with checksums) data, so that (otherwise silent) errors can be corrected
- Has been extended to deal with "fail-stop" errors
  - Maintain checksums
  - Use checksums to reconstruct lost data
- Mostly used for linear algebra applications
- Sort of like RAID 2 (parity bits, no duplication) in memory

# Reducing the checkpointing time

- One big problem is that when a checkpoint occurs, all processes say "save my state" at the same time!
  - e.g., if you have a NAS, the bandwidth to it becomes a massive bottleneck
- Several solutions have been proposed:
  - Use a cluster with fast storage (SAN)
  - Save only to local disk
    - But then you can't recover from a fatal hardware failure
  - Save to a "buddy" (save my state on my neighbor's disk)
  - Go diskless: keep my checkpoint in my buddy's memory
    - But then it must fit

# Go uncoordinated

- In many cases the "we all save at the same time" is a showstopper
- Radical idea: allow **uncoordinated** checkpointing:
  - Processes can checkpoint their data whenever they want/choose
  - Should create a "smooth/low" load on the storage system
- **Big challenge**: how do we perform recovery?
- Let's see the canonical example


# Uncoordinated checkpointing example

P2 ———

P1 ———


————→  
time

# Uncoordinated checkpointing example







# Uncoordinated checkpointing example




# Uncoordinated checkpointing example




# Uncoordinated checkpointing example




# Uncoordinated checkpointing example




# Uncoordinated checkpointing example




# Uncoordinated checkpointing example



# Uncoordinated checkpointing example



## Uncoordinated checkpointing example



Already received the blue message!!



# Message Logging

- To deal with the previous, and many more complex situations, we need to do some kind of **message logging**
- We keep track of messages sent/received, and ignore redundant copies
- How to do this efficiently and correctly is not easy
- Many "fault-tolerance MPI" implementations do this now as a matter of course
- An enormous theoretical/applied literature on this problem

# Hierarchical Checkpointing: Do it all!

- The term “hierarchical” scheduling refers loosely to a solution that applies many (all?) of the previous techniques together
- And still, projections show that this won't be enough to achieve parallel efficiency on upcoming "exascale" machines
- So we need other techniques in addition to checkpoint-rollback-recovery...

# Replication

- Run redundant processes to increase the MTBF!
  - Some MPI implementations can do this transparently!
- **Good:** we can checkpoint much less frequently
- **Bad:** we "waste" resources (and power!)
- It may seem surprising that we can gain anything by wasting resources
- But it all works out because using  $p/2$  processors has higher parallel efficiency than using  $p$  processors
- Let's see an example...

## Redundant processors can help!

- Processors fail exponentially with MTBF  $1/\lambda$
- We have an application that runs in time  $\alpha + \beta/p$  on  $p$  processors when no failures occur
- We checkpoint every  $X = \sqrt{2C/\lambda}$  (Young)
- Our makespan is:  $\alpha + \beta/p + \frac{\alpha + \beta/p}{X} \times C$
- Now we replicate each process, thus running on  $p/2$  “more reliable processors”
- The MTBF of a pair of processors is now  $(3/2)(1/\lambda)$ 
  - Compute integrals to verify this...
- We checkpoint every  $Y = \sqrt{2C(3/2)/\lambda} = \sqrt{3C/\lambda}$
- Our makespan is now:  $\alpha + 2\beta/p + \frac{\alpha + 2\beta/p}{Y} \times C$

## Redundant processors can help! (2)

- The math on the previous slide is really back-of-the-envelope (e.g., makespan values are for the failure-free case)
- Let's pick values :  $\alpha = 1$ ,  $\beta = 100$ ,  $C = 10$ ,  $\lambda = 0.001$
- For  $p = 1000$ 
  - No replication: 1.125 time units
  - Replication: 1.222 time units (replication hurts)
- For  $p = 100000$ 
  - No replication: 1.022 time units (horrendous efficiency)
  - Replication: 1.018 time units (replication helps)

# Failure Prediction

- Another idea to use in conjunction with checkpointing is **Failure Prediction**
- Hardware sensors can provide a good sense of when a processor might fail in the future
  - Based on “precursor” events
- Simple idea: checkpoint proactively only when a failure seems likely in the short term
- Raises all types of interesting questions about false positive and false negative
- This is often called “failure avoidance”

# What about silent errors?

- Dealing with silent errors has become a hot research topic
- Simple idea: perform (periodic?) checks on the data
- Trade-off:
  - Checking data infrequently leaves you open to a lot of waste
  - Checking data frequently has a lot of overhead
- Must be combined with checkpointing
- Typically we have various "checkers", some expensive and accurate, some cheap and less accurate
- Problem: deciding when and how to check for silent errors

# Fault-Tolerance for parallel applications

- This was a whirlwind tour of fault-tolerance
  - It is a huge area of research and development, that goes from pure theory to pure engineering
  - Hardware developments shape the field constantly
  - There is no current consensus on what will work for exascale, but there are many pathways
  - Every month new papers are being published :)
- 
- One recent workshop report: [\*Addressing failures in exascale computing\*](#), Snir et al., IJHPCA 28(2), 129–173, 2014.



# Outline

- 1 Why Fault Tolerance?
- 2 Checkpointing
- 3 A little theory
- 4 A little practice
- 5 Energy Consumption**

# Power is money

- Power budget is a key issue for large-scale platform
  - Hence building them next to powerplants
- Many engineering issues to reduce power consumption:
  - Microprocessor design
  - Novel cooling approaches
  - Green energies
  - ...
- But given a system with some power-management “knobs,” what do you do for a given application?
  - e.g., is your fault tolerant solution too power-hungry?
  - e.g., does your idea to decrease makespan by 1% increase power consumption by 50%?

# DVFS

- Common power-management techniques are Dynamic Voltage Scaling (DVS) and Dynamic Frequency Scaling
- Modern microprocessors allow voltage/frequency to be modified in software
- Without getting into details: you can slow down your computation and save on energy
  - e.g., Power consumption is a polynomial function of the frequency / voltage
- The question then becomes: by how much should I slow down my nodes to achieve energy and performance goals?
- Interestingly, if I take the voltage too low, then I can create silent errors (i.e., computation is wrong)

# A slew of (interesting) problems

- Say you have a parallel application and you:
  - Can run on some number of processors in a platform with DVFS-enabled processors
  - Use process duplication or not
  - Use some complex hierarchical checkpointing
  - Want to detect silent errors and have a bunch of detectors of various cost and accuracy
- How to decide on all the above?
- Sample research questions:
  - Q: minimize makespan for a given energy budget
  - Q: minimize energy given a makespan bound
  - Q: minimize chances of having a wrong result given a makespan bound and budget

## A very active research area

- Questions pertaining to fault tolerance and to energy are being investigated extremely actively
- We only scratched the surface here
- But some of the papers you'll present this semester touch on concerns of fault tolerance and power...