

# Introduction to Scheduling

## ICS632: Principles of High-Performance Computing

Henri Casanova (henric@hawaii.edu)

Fall 2015

# Foreword

- This set of lecture notes will be a bit theoretical-ish
- We'll refer to simple computational complexity concepts
- We'll have a few hand-wavy proofs
- We could have a whole semester on scheduling, this will only scratch the surface

# Outline

- 1 Scheduling Independent Tasks
- 2 Divisible Loads and OpenMP Scheduling
- 3 Scheduling Task Graphs
- 4 The Great Scheduling Zoo
- 5 Pragmatic, Dynamic Scheduling

# What is scheduling?

- Broad definition: *the temporal allocation of activities to resources to achieve some desirable objective*
- Examples:
  - Assign workers to machines in an factory to increase productivity
  - Pick classrooms for classes at a university to maximize the number of free classrooms on Fridays
  - Assign users to a pay-per-hour telescope to maximize profit
  - **Assign computation to processors and communications to network links so as to minimize application execution time**

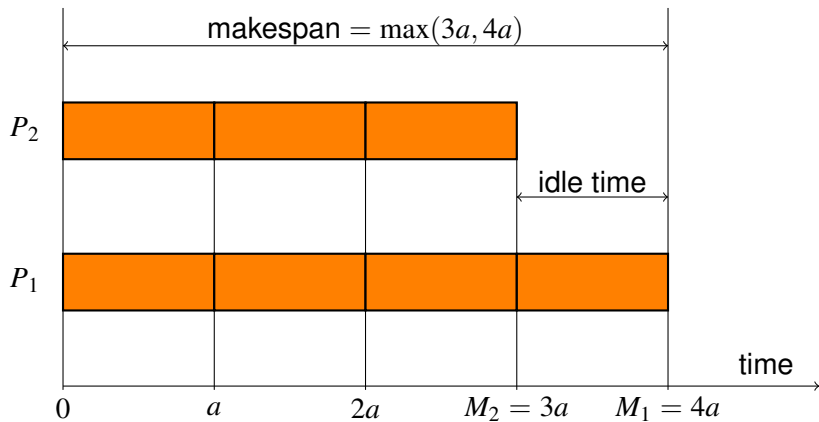
# A simple scheduling problem

- A Scheduling Problem is defined by three components:
  - 1 A description of a set of resources
  - 2 A description of a set of tasks
  - 3 A description of a desired objective
- Let us get started with a simple problem: INDEP(2)
  - 1 Two identical processors,  $P_1$  and  $P_2$ 
    - Each processor can run only one task at a time
  - 2  $n$  compute tasks
    - Each task can run on either processor in  $a$  seconds
    - Tasks are *independent*: can be computed in any order
  - 3 Objective: minimize  $\max(M_1, M_2)$  (**makespan**)
    - $M_i$  is the time at which processor  $P_i$  finishes computing

## The easy case

- If all tasks are *identical*, i.e., take the same amount of compute time, then the solution is obvious: Assign  $\lceil n/2 \rceil$  tasks to  $P_1$  and  $\lfloor n/2 \rfloor$  tasks to  $P_2$ 
  - Rule of thumb: try to have both processors finish at the same time
- We have a trivial linear-time algorithm
  - For each task pick one of the two processors by comparing the index of the task with  $n/2$
- In fact we have already seen an optimal algorithm for a more complex situation in which we have  $p$  *heterogeneous* processors

# Gantt chart for INDEP(2) with 7 identical tasks



# Non-identical tasks

- Task  $T_i$ ,  $i = 1, \dots, n$  takes time  $a_i \geq 0$
- There is no p-time algorithm to solve INDEP(2) (unless  $\mathcal{P} = \mathcal{NP}$ )
- INDEP(2) (*decision* version) is in  $\mathcal{NP}$ 
  - Certificate: for each  $a_i$  whether it is scheduled on  $P_1$  or  $P_2$
  - In linear time, compute the makespan on both processors, and compare to makespan bound to answer “Yes”
- Consider an instance of 2-PARTITION ( $\mathcal{NP}$ -complete):
  - Given  $n$  integers  $x_i$ , is there a subset  $I$  of  $\{1, \dots, n\}$  such that  $\sum_{i \in I} x_i = \sum_{i \notin I} x_i$ ?
- Let us construct an instance of INDEP(2):
  - Let  $k = \frac{1}{2} \sum x_i$ , let  $a_i = x_i$
- The proof is trivial
  - If  $k$  is non-integer, neither instance has a solution
  - Otherwise, each processor corresponds to one subset
- INDEP(2) is identical to 2-PARTITION



# So what?

- This  $\mathcal{NP}$ -completeness proof is probably the most trivial in the world 😊
- But now we are thus pretty sure that there is no p-time algorithm to solve INDEP(2)
- What we look for now are *approximation algorithms...*

# Approximation algorithms

- Consider an optimization problem
- A  $p$ -time algorithm is a  $\lambda$ -*approximation algorithm* if it returns a solution that's at most a factor  $\lambda$  from the optimal solution (the closer  $\lambda$  to 1, the better)
  - $\lambda$  is called the *approximation ratio*
- *Polynomial Time Approximation Scheme* (PTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm (may be non-polynomial is  $1/\epsilon$ )
- *Fully Polynomial Time Approximation Scheme* (FPTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm polynomial in  $1/\epsilon$
- Typical goal: find a FPTAS, if not find a PTAS, if not find a  $\lambda$ -approximation for a low value of  $\lambda$

# Approximation algorithms

- Consider an optimization problem
- A  $p$ -time algorithm is a  $\lambda$ -*approximation algorithm* if it returns a solution that's at most a factor  $\lambda$  from the optimal solution (the closer  $\lambda$  to 1, the better)
  - $\lambda$  is called the *approximation ratio*
- *Polynomial Time Approximation Scheme* (PTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm (may be non-polynomial in  $1/\epsilon$ )
- *Fully Polynomial Time Approximation Scheme* (FPTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm polynomial in  $1/\epsilon$
- Typical goal: find a FPTAS, if not find a PTAS, if not find a  $\lambda$ -approximation for a low value of  $\lambda$

# Approximation algorithms

- Consider an optimization problem
- A  $p$ -time algorithm is a  $\lambda$ -*approximation algorithm* if it returns a solution that's at most a factor  $\lambda$  from the optimal solution (the closer  $\lambda$  to 1, the better)
  - $\lambda$  is called the *approximation ratio*
- *Polynomial Time Approximation Scheme* (PTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm (may be non-polynomial is  $1/\epsilon$ )
- *Fully Polynomial Time Approximation Scheme* (FPTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm polynomial in  $1/\epsilon$
- Typical goal: find a FPTAS, if not find a PTAS, if not find a  $\lambda$ -approximation for a low value of  $\lambda$

# Approximation algorithms

- Consider an optimization problem
- A  $p$ -time algorithm is a  $\lambda$ -*approximation algorithm* if it returns a solution that's at most a factor  $\lambda$  from the optimal solution (the closer  $\lambda$  to 1, the better)
  - $\lambda$  is called the *approximation ratio*
- *Polynomial Time Approximation Scheme* (PTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm (may be non-polynomial is  $1/\epsilon$ )
- *Fully Polynomial Time Approximation Scheme* (FPTAS): for any  $\epsilon$  there exists a  $(1 + \epsilon)$ -approximation algorithm polynomial in  $1/\epsilon$
- Typical goal: find a FPTAS, if not find a PTAS, if not find a  $\lambda$ -approximation for a low value of  $\lambda$


# Greedy algorithms

- A greedy algorithm is one that builds a solution step-by-step, via local incremental decisions
- It turns out that several greedy scheduling algorithms are approximation algorithms
  - Informally, they're not as “bad” as one may think
- Two natural greedy algorithms for INDEP(2):
  - **greedy-online**: take the tasks in arbitrary order and assign each task to the least loaded processor
    - As if we don't know which tasks are coming
  - **greedy-offline**: sort the tasks by decreasing  $a_i$ , and assign each task in that order to the least loaded processor
    - We know all the tasks ahead of time

# Greedy algorithms



- A greedy algorithm is one that builds a solution step-by-step, via local incremental decisions
- It turns out that several greedy scheduling algorithms are approximation algorithms
  - Informally, they're not as “bad” as one may think
- Two natural greedy algorithms for INDEP(2):
  - **greedy-online**: take the tasks in arbitrary order and assign each task to the least loaded processor
    - As if we don't know which tasks are coming
  - **greedy-offline**: sort the tasks by decreasing  $a_i$ , and assign each task in that order to the least loaded processor
    - We know all the tasks ahead of time

## Example with 6 tasks: Online

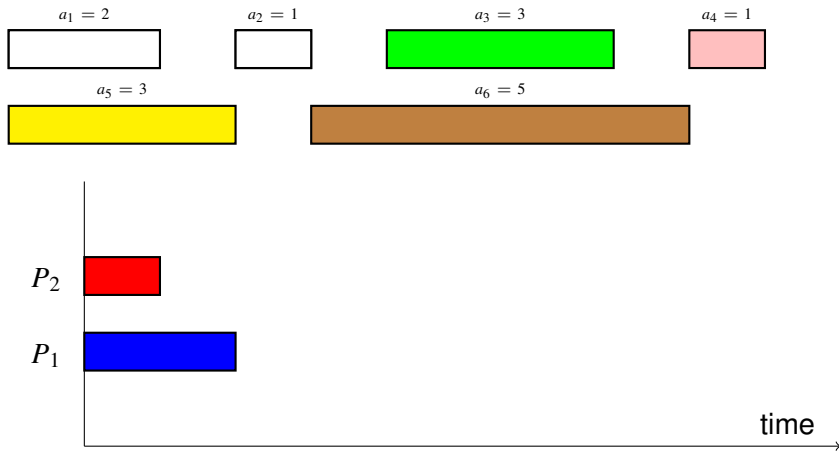




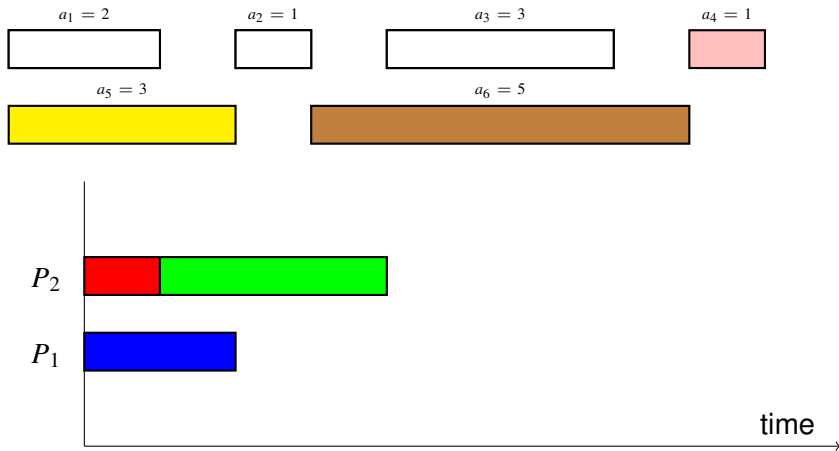
# Example with 6 tasks: Online





# Example with 6 tasks: Online





## Example with 6 tasks: Online





## Example with 6 tasks: Online





## Example with 6 tasks: Online




## Example with 6 tasks: Online



## Example with 6 tasks: Online





# Example with 6 tasks: Offline








## Example with 6 tasks: Offline





## Example with 6 tasks: Offline





## Example with 6 tasks: Offline





## Example with 6 tasks: Offline





## Example with 6 tasks: Offline



## Example with 6 tasks: Offline



## Example with 6 tasks: Offline



# Greedy-online for INDEP(2)

## Theorem

*Greedy-online is a  $\frac{3}{2}$ -approximation*

### ■ Proof:

- $P_i$  finishes computing at time  $M_i$  ( $M$  stands for makespan)
- Let us assume  $M_1 \geq M_2$  ( $M_{greedy} = M_1$ )
- Let  $T_j$  the last task to execute on  $P_1$
- Since the greedy algorithm put  $T_j$  on  $P_1$ , then  $M_1 - a_j \leq M_2$
- We have  $M_1 + M_2 = \sum_i a_i = S$
- $M_{greedy} = M_1 = \frac{1}{2}(M_1 + (M_1 - a_j) + a_j) \leq \frac{1}{2}(M_1 + M_2 + a_j) = \frac{1}{2}(S + a_j)$
- but  $M_{opt} \geq S/2$  (ideal lower bound on optimal)
- and  $M_{opt} \geq a_j$  (at least one task is executed)
- Therefore:  $M_{greedy} \leq \frac{1}{2}(2M_{opt} + M_{opt}) = \frac{3}{2}M_{opt}$   $\square$



# Greedy-online for INDEP(2)

## Theorem

*Greedy-online is a  $\frac{3}{2}$ -approximation*

### ■ Proof:

- $P_i$  finishes computing at time  $M_i$  ( $M$  stands for makespan)
- Let us assume  $M_1 \geq M_2$  ( $M_{greedy} = M_1$ )
- Let  $T_j$  the last task to execute on  $P_1$
- Since the greedy algorithm put  $T_j$  on  $P_1$ , then  $M_1 - a_j \leq M_2$
- We have  $M_1 + M_2 = \sum_i a_i = S$
- $M_{greedy} = M_1 = \frac{1}{2}(M_1 + (M_1 - a_j) + a_j) \leq \frac{1}{2}(M_1 + M_2 + a_j) = \frac{1}{2}(S + a_j)$
- but  $M_{opt} \geq S/2$  (ideal lower bound on optimal)
- and  $M_{opt} \geq a_j$  (at least one task is executed)
- Therefore:  $M_{greedy} \leq \frac{1}{2}(2M_{opt} + M_{opt}) = \frac{3}{2}M_{opt}$   $\square$

# Greedy-online for INDEP(2)

## Theorem

*Greedy-online is a  $\frac{3}{2}$ -approximation*

### ■ Proof:

- $P_i$  finishes computing at time  $M_i$  ( $M$  stands for makespan)
- Let us assume  $M_1 \geq M_2$  ( $M_{\text{greedy}} = M_1$ )
- Let  $T_j$  the last task to execute on  $P_1$
- Since the greedy algorithm put  $T_j$  on  $P_1$ , then  $M_1 - a_j \leq M_2$
- We have  $M_1 + M_2 = \sum_i a_i = S$
- $M_{\text{greedy}} = M_1 = \frac{1}{2}(M_1 + (M_1 - a_j) + a_j) \leq \frac{1}{2}(M_1 + M_2 + a_j) = \frac{1}{2}(S + a_j)$
- but  $M_{\text{opt}} \geq S/2$  (ideal lower bound on optimal)
- and  $M_{\text{opt}} \geq a_j$  (at least one task is executed)
- Therefore:  $M_{\text{greedy}} \leq \frac{1}{2}(2M_{\text{opt}} + M_{\text{opt}}) = \frac{3}{2}M_{\text{opt}} \quad \square$

## Greedy-offline for INDEP(2)

## Theorem

*Greedy-offline is a  $\frac{7}{6}$ -approximation*

■ Proof:

- If  $a_j \leq \frac{1}{3}M_{opt}$ , the previous proof can be used
  - $M_{greedy} \leq \frac{1}{2}(2M_{opt} + \frac{1}{3}M_{opt}) = \frac{7}{6}M_{opt}$
- If  $a_j > \frac{1}{3}M_{opt}$ , then  $j \leq 4$ 
  - if  $T_j$  was the 5th task, then, due to the task ordering, there would be 5 tasks with  $a_i > \frac{1}{3}M_{opt}$
  - There would be at least 3 tasks on the same processor in the optimal schedule
  - Therefore  $M_{opt} > 3 \times \frac{1}{3}M_{opt}$ , a contradiction
- One can check all possible scenarios for 4 tasks and show optimality □

# Greedy-offline for INDEP(2)

## Theorem

*Greedy-offline is a  $\frac{7}{6}$ -approximation*

### ■ Proof:

- If  $a_j \leq \frac{1}{3}M_{opt}$ , the previous proof can be used
  - $M_{greedy} \leq \frac{1}{2}(2M_{opt} + \frac{1}{3}M_{opt}) = \frac{7}{6}M_{opt}$
- If  $a_j > \frac{1}{3}M_{opt}$ , then  $j \leq 4$ 
  - if  $T_j$  was the 5th task, then, due to the task ordering, there would be 5 tasks with  $a_i > \frac{1}{3}M_{opt}$
  - There would be at least 3 tasks on the same processor in the optimal schedule
  - Therefore  $M_{opt} > 3 \times \frac{1}{3}M_{opt}$ , a contradiction
- One can check all possible scenarios for 4 tasks and show optimality    □

# Bounds are tight

## ■ Greedy-online:

- $a_i$ 's =  $\{1, 1, 2\}$
- $M_{greedy} = 3; M_{opt} = 2$
- $ratio = \frac{3}{2}$

## ■ Greedy-offline:

- $a_i$ 's =  $\{3, 3, 2, 2, 2\}$
- $M_{greedy} = 7; M_{opt} = 6$
- $ratio = \frac{7}{6}$

# PTAS and FPTAS for INDEP(2)

## Theorem

*There is a PTAS  $((1 + \epsilon)$ -approximation) for INDEP(2)*

### ■ Proof Sketch:

- Classify tasks as either “small” or “large”
  - Very common technique
- Replace all small tasks by same-size tasks
- Compute an optimal schedule of the modified problem in p-time (not polynomial in  $1/\epsilon$ )
- Show that the cost is  $\leq 1 + \epsilon$  away from the optimal cost
- The proof is a couple of pages, but not terribly difficult

## Theorem

*There is a FPTAS  $((1 + \epsilon)$ -approx pol. in  $1/\epsilon$ ) for INDEP(2)*

## We know a lot about INDEP(2)

- INDEP(2) is NP-complete
  - We have simple greedy algorithms with guarantees on result quality
  - We have a simple PTAS
  - We even have a (less simple) FPTAS
  - INDEP(2) is basically “solved”
- 
- Sadly, not many scheduling problems are this well-understood...

# INDEP(P) is much harder

- INDEP(P) is  $\mathcal{NP}$ -complete by trivial reduction to 3-PARTITION:
  - Give  $3n$  integers  $a_1, \dots, a_{3n}$  and an integer  $B$ , can we partition the  $3n$  integers into  $n$  sets, each of sum  $B$ ?  
(assuming that  $\sum_i a_i = nB$ )
- 3-PARTITION is  $\mathcal{NP}$ -complete “in the strong sense”, unlike 2-PARTITION
  - Even when encoding the input in unary (i.e., no logarithmic numbers of bits), one cannot find an algorithm polynomial in the size of the input!
  - Informally, a problem is  $\mathcal{NP}$ -complete “in the weak sense” if it is hard only if the numbers in the input are unbounded
- INDEP(P) is thus fundamentally harder than INDEP(2)



# Approximation algorithm for INDEP(P)

## Theorem

*Greedy-online is a  $(2 - \frac{1}{p})$ -approximation*

### ■ Proof (usual reasoning):

- Let  $M_{greedy} = \max_{1 \leq i \leq p} M_i$ , and  $j$  be such that  $M_j = M_{greedy}$
- Let  $T_k$  be the last task assigned to processor  $P_j$
- $\forall i, \quad M_i \geq M_j - a_k$  (greedy algorithm)
- $S = \sum_i^p M_i = M_j + \sum_{i \neq j} M_i \geq M_j + (p-1)(M_j - a_k) = pM_j + (p-1)a_k$
- Therefore,  $M_{greedy} = M_j \leq \frac{S}{p} + (1 - \frac{1}{p})a_k$
- But  $M_{opt} \geq a_k$  and  $M_{opt} \geq S/p$
- So  $M_{greedy} \leq M_{opt} + (1 - \frac{1}{p})M_{opt} \quad \square$
- This ratio is “tight” (e.g., an instance with  $p(p-1)$  tasks of size 1 and one task of size  $p$  has this ratio)

# Approximation algorithm for INDEP(P)

## Theorem

*Greedy-offline is a  $(\frac{4}{3} - \frac{1}{3p})$ -approximation*

- The proof is more involved, but follows the spirit of the proof for INDEP(2)
- This ratio is tight
- There is a PTAS for INDEP(P), a  $(1 + \epsilon)$ -approximation (massively exponential in  $1/\epsilon$ )
- There is no known FPTAS, unlike for INDEP(2)

# Approximation algorithm for INDEP(P)

## Theorem

*Greedy-offline is a  $(\frac{4}{3} - \frac{1}{3p})$ -approximation*

- The proof is more involved, but follows the spirit of the proof for INDEP(2)
- This ratio is tight
- There is a PTAS for INDEP(P), a  $(1 + \epsilon)$ -approximation (massively exponential in  $1/\epsilon$ )
- There is no known FPTAS, unlike for INDEP(2)


# Outline

- 1 Scheduling Independent Tasks
- 2 Divisible Loads and OpenMP Scheduling**
- 3 Scheduling Task Graphs
- 4 The Great Scheduling Zoo
- 5 Pragmatic, Dynamic Scheduling

# Why are many scheduling problems hard?


- Many scheduling problems are  $\mathcal{NP}$ -complete
- One contributing reason is that they involve integer constraints
  - The same reason why bin packing is difficult: you can't cut boxes into pieces!
- Let's see this on an example...

## INDEP(P) example schedule (offline)




$$\sum a_i = 21; \text{makespan} = 8$$

# INDEP(P) example schedule (offline)




Let's modify the schedule using preemption/migration

# INDEP(P) example schedule (offline)






# INDEP(P) example schedule (offline)



# INDEP(P) example schedule (offline)



$$\sum a_i = 21; \text{ makespan} = 7 \text{ (optimal: no idle time)}$$

# Cutting tasks

- By “cutting” a task in two, we’re able to have all processors finish at the same time
  - Zero idle time means the schedule is optimal
- If we were able to cut all tasks into tiny bits, then we would always be able to achieve zero idle time
  - Again, if you have a knife, binpacking is easy
  - Of course, there’d be “cutting overhead”...
- Question: Can this be done for real-world applications?

# Divisible Load applications

- It turns out that many useful applications consist of very large numbers of small, independent, and identical tasks
  - task execution time  $\ll$  application execution time
  - tasks can be completed in any order
  - tasks all do the same thing, but on different data
- Example applications:
  - Ray tracing (1 task = 1 photon)
  - MPEG encoding of a movie (1 task = 1 frame)
  - Seismic event processing (1 task = 1 event)
  - High-Energy Physics (1 task = 1 particle)
- These applications are termed *Divisible Loads* (DLs)
  - So fine-grain that a *continuous load* assumption is valid
- This should make scheduling trivial (INDEP(P) with same-size tasks)

# OpenMP Loops

- OpenMP is used primarily to parallelize loops in which all iterations are independent
- If the number of iterations is large, a loop is a divisible load!
- Simple divisible load assumption: If  $n$  iterations on  $p$  cores, then each core performs  $\sim n/p$  iterations
- Easy, right?
- But:
  - 1 Not all iterations are always equal
    - So we want to create a lot of chunks to avoid idle time!
  - 2 Creating a chunk of iterations incurs overhead
    - So we want to create few chunks to avoid overhead!

# OpenMP: chunk size $\sim n/p$

## OpenMP

```
#pragma omp parallel for schedule(static)
for (i=0; i < N; i++) {
    // compute something
}
```

- Each thread performs  $\sim n/p$  iterations
- Low overhead: “assign” work to each thread once
- High potential idle time if iterations are non-identical

# OpenMP: chunk size = constant

## OpenMP

```
#pragma omp parallel for schedule(dynamic, chunksize)
for (i=0; i < N; i++) {
    // compute something
}
```

- Each thread performs *chunksize* iterations (default = 1)
- High overhead (if  $chunksize \ll N$ ): “assign” work to each thread many times
  - Implemented via a critical section to increment an index
- Low idle time (if  $chunksize \ll N$ ): if iterations are wildly different then using  $chunksize = 1$  corresponds to the on-line optimal algorithm for solving INDEP(P), but with overhead added to each task

# OpenMP: chunk size = variable

## OpenMP

```
#pragma omp parallel for schedule(guided, min_chunksize)
for (i=0; i < N; i++) {
    // compute something
}
```

- Chunk sizes are created as follows and executed in an greedy fashion in this order
  - $N/2$  iterations partitioned into  $p$  chunks
  - $N/4$  iterations partitioned into  $p$  chunks
  - ...
  - until a minimal chunksize is reached (default=1)
- Goal:
  - Low overhead at the beginning, no idle time anyway
  - High overhead at the end but low idle time




# Outline

- 1 Scheduling Independent Tasks
- 2 Divisible Loads and OpenMP Scheduling
- 3 Scheduling Task Graphs**
- 4 The Great Scheduling Zoo
- 5 Pragmatic, Dynamic Scheduling

# Task dependencies

- In practice tasks often have *dependencies*
- A general model of computation is the Acyclic Directed Graph (DAG),  $G = (V, E)$
- Each task has a *weight* (i.e., execution time in seconds), a *parent*, and *children*
- The first task is the *source*, the last task the *sink*
- Topological (partial) order of the tasks





# Where do DAGs come from?

- Consider a (lower) triangular linear system solve
  - What you would need to do after an LU factorization

$$Ax = b \quad \begin{array}{|c|} \hline \text{[Lower Triangular Matrix]} \\ \hline \end{array} * \begin{array}{|c|} \hline \text{[Vector b]} \\ \hline \end{array} = \begin{array}{|c|} \hline \text{[Vector x]} \\ \hline \end{array}$$

- Simple Algorithm

```
for (i = 0; i < n; i++) {  
    x[i] = b[i] / a[i,i];  
    for (j=i+1; j<n; j++) {  
        b[j] = b[j] - a[j,i] * x[i];  
    }  
}
```



# Where do DAGs come from?

- Consider a (lower) triangular linear system solve
  - What you would need to do after an LU factorization

$$Ax = b$$


- Simple Algorithm

```
for (i = 0; i < n; i++) {  
    Ti,i: x[i] = b[i] / a[i,i];  
    for (j=i+1; j<n; j++) {  
        Ti,j: b[j] = b[j] - a[j,i] * x[i];  
    }  
}
```



# Tasks, Dependencies, etc.

```
for (i = 0; i < n; i++) {  
    Ti,*: x[i] = b[i] / a[i,i];  
    for (j=i+1; j<n; j++) {  
        Ti,j: b[j] = b[j] - a[j,i] * x[i];  
    }  
}
```

- All tasks  $T_{i,*}$  are executed at iteration  $i$  of the outer loop

- There is a simple sequential order of the tasks

$$T_{0,0} < T_{0,1} < \dots < T_{0,n-1} < T_{1,0} < T_{1,1} < \dots < T_{1,n-1} < \dots$$

- Of course, when considering a parallel execution, one tries to find **independent** tasks
- To see if tasks are independent one must examine their input (In) and their output (Out)



# Tasks, Dependencies, etc.

```
for (i = 0; i < n; i++) {  
    Ti,i: x[i] = b[i] / a[i,i];  
    for (j=i+1; j<n; j++) {  
        Ti,j: b[j] = b[i] - a[j,i] * x[i];  
    }  
}
```

## ■ Input and Output

- $\text{In}(T_{i,i}) = \{b[i], a[i,i]\}$
- $\text{Out}(T_{i,i}) = \{x[i]\}$
- $\text{In}(T_{i,j}) = \{b[i], a[j,i], x[i]\}$  for  $j > i$
- $\text{Out}(T_{i,j}) = \{b[j]\}$  for  $j > i$

## ■ Bernstein Conditions

- T and T' are independent if all 3 conditions are met
  - $\text{In}(T) \cap \text{Out}(T') = \emptyset$
  - $\text{Out}(T) \cap \text{In}(T') = \emptyset$
  - $\text{Out}(T) \cap \text{Out}(T') = \emptyset$



# Task Graph


```
for (i = 0; i < n; i++) {  
     $T_{i,i}$ :  $x[i] = b[i] / a[i,i];$   
    for (j=i+1; i<n; i++) {  
         $T_{i,j}$ :  $b[j] = b[j] - a[j,i] * x[i];$   
    }  
}
```

- It is easy to see that
  - for all  $i$ , all  $T_{i,j}$  are independent of each other for  $j > i$
  - for all  $i$ , all  $T_{i,j}$  depend on  $T_{i,i}$ , for  $j > i$
  - for all  $i$ , all  $T_{i,j}$  depend on  $T_{i-1,j}$  for  $j \geq i$  and  $i > 0$
- Hence the task graph



# Task Graph

```
for (i = 0; i < n; i++) {  
  Ti,i: x[i] = b[i] / a[i,i];  
  for (j=i+1; i<n; i++) {  
    Ti,j: b[j] = b[i] - a[j,i] * x[i];  
  }  
}
```








# More taskgraphs

- The previous taskgraph comes from a low-level analysis of the code
  - It probably makes little sense to do a parallel implementation with MPI with such a low task granularity
  - Can totally make sense with OpenMP
  - Such task graphs can also be used by compilers to do code optimization by exploiting multiple functional units, pipelines functional units, etc.
  - With “blocking” these tasks could become MPI tasks
- Other taskgraphs are really how the application was build



# Scientific Workflows


- A popular way in which many scientific applications are constructed is as **workflows**
  - A scientist conceptually drags and drops computational kernels and connects their input-output
  - The result is a DAG (actually more general than a DAG) that does something useful
- Example Application: Montage
  - Produce Mosaic of the Sky
  - Based on multiple data sources
  - Given angle, coordinates, size, etc.
  - 10s of thousands of tasks
- Example: M101 galaxy images





# Many levels of parallelisms


- Montage Workflow





# Many levels of parallelisms


- Montage Workflow






# Many levels of parallelisms

- Montage workflow



# Critical path

- Assume that the DAG executes on  $p$  processors
- The longest path (in seconds) is called the *critical path*
- The length of the critical path (CP) is a *lower bound on  $M_{opt}$* , regardless of the number of processors
- In this example, the CP length is 6 (the other path has length 4)



# Complexity

- Unsurprisingly, DAG scheduling is  $\mathcal{NP}$ -complete
  - Independent tasks is a special case of DAG scheduling
- Typical greedy algorithm skeleton:
  - Maintain a list of *ready* tasks (with cleared dependencies)
  - Greedily assign a ready task to an available processor as early as possible (don't leave a processor idle unnecessarily)
  - Update the list of ready tasks
  - Repeat until all tasks have been scheduled
- This is called **List Scheduling**
- Many list scheduling algorithms are possible
  - Depending on how to select the ready task to schedule next


# Complexity

- Unsurprisingly, DAG scheduling is  $\mathcal{NP}$ -complete
  - Independent tasks is a special case of DAG scheduling
- Typical greedy algorithm skeleton:
  - Maintain a list of *ready* tasks (with cleared dependencies)
  - Greedily assign a ready task to an available processor as early as possible (don't leave a processor idle unnecessarily)
  - Update the list of ready tasks
  - Repeat until all tasks have been scheduled
- This is called **List Scheduling**
- Many list scheduling algorithms are possible
  - Depending on how to select the ready task to schedule next




# List scheduling example

3 Processors




# List scheduling example



3 Processors




# List scheduling example



3 Processors




# List scheduling example



3 Processors




# List scheduling example




3 Processors




# List scheduling example




3 Processors




# List scheduling example




3 Processors



# List scheduling example




3 Processors






# List scheduling example



3 Processors



Makespan = 16; CP Length = 15

Idle Time = 1+5+5+8 = 19

# List scheduling

## Theorem (fundamental)

*List scheduling is a  $(2 - \frac{1}{p})$ -approximation*

- Doesn't matter how the next ready task is selected
- Let's prove this theorem informally
  - Really simple proof if one doesn't use the typical notations for schedules
  - I never use these notations in public 😊

# List scheduling


## Theorem (fundamental)

*List scheduling is a  $(2 - \frac{1}{p})$ -approximation*

- Doesn't matter how the next ready task is selected
- Let's prove this theorem informally
  - Really simple proof if one doesn't use the typical notations for schedules
  - I never use these notations in public 😊


# Approximation ratio

Let's consider a list-scheduling schedule



# Approximation ratio


Let's consider one of the tasks that finishes last



# Approximation ratio


Why didn't this task run during an earlier idle period?

Because a parent was not finished (list scheduling!)



# Approximation ratio


Let's look at a parent



# Approximation ratio

Why didn't this task run during an earlier idle period?


Because a parent was not finished (list scheduling!)





# Approximation ratio

Let's look at a parent



## Approximation ratio


## Why didn't this task run during an earlier idle period?

Because a parent was not finished (list scheduling!)




# Approximation ratio

Let's look at a parent




# Approximation ratio

And so on...



# Approximation ratio

At any point in time either a task on the red path is running or no processor is idle



# Approximation ratio

- Let  $L$  be the length of the red path (in seconds),  $p$  the number of processors,  $I$  the total idle time,  $M$  the makespan, and  $S$  the sum of all task weights
  - $I \leq (p - 1)L$ 
    - processors can be idle only when a red task is running
  - $L \leq M_{opt}$ 
    - The optimal makespan is longer than any path in the DAG
  - $M_{opt} \geq S/p$ 
    - $S/p$  is the makespan with zero idle time
  - $p \times M = I + S$ 
    - rectangle's area = white boxes + non-white boxes
- $\Rightarrow p \times M \leq (p - 1)M_{opt} + pM_{opt} \Rightarrow M \leq (2 - \frac{1}{p})M_{opt} \quad \square$

# Good list scheduling?

- All list scheduling algorithms thus have the same approximation ratio
- But there are many options for list scheduling
  - Many ways of sorting the ready tasks...
- In practice, some may be better than others
- One well-known option, *Critical path scheduling*

# Critical path scheduling

- When given a set of ready tasks, which one do we pick to schedule?
- Idea: pick a task on the CP
  - If we prioritize tasks on the CP, then the CP length is reduced
  - The CP length is a lower bound on the makespan
  - So intuitively it's good for it to be low
- For each (ready) task, compute its *bottom level*, the length of the path from the task to the sink
- Pick the task with the *largest* bottom level



# Outline

- 1 Scheduling Independent Tasks
- 2 Divisible Loads and OpenMP Scheduling
- 3 Scheduling Task Graphs
- 4 The Great Scheduling Zoo**
- 5 Pragmatic, Dynamic Scheduling

# Graham's notation

- There are SO many variations on the scheduling problem that Graham has proposed a standard notation:  $\alpha|\beta|\gamma$ 
  - *alpha*: processors
  - *beta*: tasks
  - *gamma*: objective function
- Let's see some examples for each

## $\alpha$ : processors

- 1: one processor
- $Pn$ :  $n$  identical processors (if  $n$  not fixed, not given)
- $Qn$ :  $n$  uniform processors (if  $n$  not fixed, not given)
  - Each processor has a (different) compute speed
- $Rn$ :  $n$  unrelated processors (if  $n$  not fixed, not given)
  - Each processor has a (different) compute speed for each (different) task (e.g.,  $P_1$  can be faster than  $P_2$  for  $T_1$ , but slower for  $T_2$ )

## $\beta$ : tasks

- $r_j$ : tasks have *release dates*
- $d_j$ : tasks have *deadlines*
- $p_j = x$ : all tasks have weight  $x$
- $prec$ : general precedence constraints (DAG)
- $tree$ : tree precedence constraints
- $chains$ : chains precedence constraints (multiple independent paths)
- $pmtn$ : tasks can be preempted and restarted (on other processors)
  - Makes scheduling easier, and can often be done in practice
- ...

## $\gamma$ : objective function

- $C_{max}$ : makespan
- $\sum C_i$ : mean flow-time (completion time minus release date if any)
- $\sum w_i C_i$ : average weighted flow-time
- $L_{max}$ : maximum lateness ( $\max(0, C_i - d_i)$ )
- ...

# Example scheduling problems

- The classification is not perfect and variations among authors are common
- Some examples:
  - $P2||C_{max}$ , which we called INDEP(2)
  - $P||C_{max}$ , which we called INDEP(P)
  - $P|prec|C_{max}$ , which we called DAG scheduling
  - $R2|chains|\sum C_i$ 
    - Two related processors, chains, minimize sum-flow
  - $P|r_j; p_j \in \{1, 2\}; d_j; pmtn|L_{max}$ 
    - Identical processors, tasks with release dates and deadlines, task weights either 1 or 2, preemption, minimize maximum lateness

# Where to find known results

- Luckily, the body of knowledge is well-documented (and Graham's notation widely used)
- Several books on scheduling that list known results
  - *Handbook of Scheduling*, Leung and Anderson
  - *Scheduling Algorithms*, Brucker
  - *Scheduling: Theory, Algorithms, and Systems*, Pinedo
  - ...
- Many published survey articles

# Example list of known results

■ Excerpt from  
*Scheduling  
Algorithm*, P.  
Brucker

$P2 \parallel C_{max}$	Lenstra et al. [155]
* $P \parallel C_{max}$	Garey & Johnson [98]
* $P \mid p_i = 1;intree;r_i \mid C_{max}$	Brucker et al. [35]
* $P \mid p_i = 1;prec \mid C_{max}$	Ullman [203]
* $P2 \mid chains \mid C_{max}$	Du et al. [86]
* $Q \mid p_i = 1;chains \mid C_{max}$	Kubiak [129]
* $P \mid p_i = 1;outtree \mid L_{max}$	Brucker et al. [35]
* $P \mid p_i = 1;intree;r_i \mid \sum C_i$	Lenstra [150]
* $P \mid p_i = 1;prec \mid \sum C_i$	Lenstra & Rinnooy Kan [152]
* $P2 \mid chains \mid \sum C_i$	Du et al. [86]
* $P2 \mid r_i \mid \sum C_i$	Single-machine problem
$P2 \parallel \sum w_i C_i$	Bruno et al. [58]
* $P \parallel \sum w_i C_i$	Lenstra [150]
* $P2 \mid p_i = 1;chains \mid \sum w_i C_i$	Timkovsky [201]
* $P2 \mid p_i = 1;chains \mid \sum U_i$	Single-machine problem
* $P2 \mid p_i = 1;chains \mid \sum T_i$	Single-machine problem

Table 5.3:  $\mathcal{NP}$ -hard parallel machine problems without preemption.



# Outline

- 1 Scheduling Independent Tasks
- 2 Divisible Loads and OpenMP Scheduling
- 3 Scheduling Task Graphs
- 4 The Great Scheduling Zoo
- 5 Pragmatic, Dynamic Scheduling**

# Does any of this help?

- So far we've looked at very simple models
- What if we throw everything in?
  - Non-identical tasks
  - Task dependencies
  - Heterogeneous compute nodes
  - Complex network topologies with heterogeneous “links”
  - Task computation times not known ahead of time
  - Tasks not known ahead of time
  - Tasks are themselves parallel
  - Processor/network speeds change
  - etc.
- But how do we even reason about something this complicated and non-deterministic?
- One good practical option: dynamic execution

# Master-Worker Dynamic execution

- Master process:
  - Keeps a list of “ready” tasks to compute with where input data can be found
- Worker processes
  - Request work from the master **when idle**
  - Create new tasks to compute and tell the master about them
- Each work request causes overhead
- Cleverness should be used to avoid long data transfers
  - i.e., leave data distributed and have the master try to assign a task to a process that already has the data


# Work Stealing

- No Master
- Each process keeps a queue of tasks to perform and extracts a task from its local queue whenever idle
- If the queue is empty, then a process *steals* a task from another “victim” process
- Many options:
  - How many candidate victims should be considered?
  - Which victim do I pick?
  - How many tasks to I steal (overhead)?
- Under some assumptions, there are theoretical results:
  - Bounds on numbers of steals, with high probability
  - Bounds on overall makespan, with high probability
- Many efficient implementations in shared-memory (e.g., Cilk) and distributed-memory (e.g., Kaapi)


# Conclusion

- Scheduling problems are diverse and often difficult
- Relevant theoretical questions:
  - Is it in  $\mathcal{P}$ ?
  - Is it  $\mathcal{NP}$ -complete?
    - Are there approximation algorithms?
    - Are there PTAS or FTPAS?
    - Are there at least decent non-guaranteed heuristics?
- Luckily, scheduling problems have been studied a lot
- Come up with the Graham notation for your problem and check what is known about it!
- In the wild, dynamic scheduling may work well


# Sources



Y. Robert  
F. Vivien



H. Casanova  
A. Legrand  
Y. Robert



A. Benoit  
Y. Robert  
F. Vivien