Simulation

ICS632: Principles of High-Performance Computing

Henri Casanova (henric@hawaii.edu)

Fall 2015

Acknowledgments

Joint work with a LOT of people:

http://simgrid.gforge.inria.fr



Why Simulation?

- Many "HPC" theoretical results
 - Optimal alg., approx. alg., \mathcal{NP} -completeness
- Theory is key to understanding problems better
- Theoretical results also provide motivation and inspiration for developing non-guaranteed heuristics
 - e.g., when there is no optimal algorithm
 - e.g., non-guaranteed algorithms may achieve much better average perfor mance than approximation algorithms
- Sadly, we rarely have theorems to compare heuristics
- Typically, given 20 (reasonable) heuristics and 10,000 random problem instance s, each heuristic could be best on some of the instances
- Question: How do we find out which heuristics are best?



Why Simulation?

- Many "HPC" theoretical results
 - Optimal alg., approx. alg., NP-completeness
- Theory is key to understanding problems better
- Theoretical results also provide motivation and inspiration for developing non-guaranteed heuristics
 - e.g., when there is no optimal algorithm
 - e.g., non-guaranteed algorithms may achieve much better average perfor mance than approximation algorithms
- Sadly, we rarely have theorems to compare heuristics
- Typically, given 20 (reasonable) heuristics and 10,000 random problem instance s, each heuristic could be best on some of the instances
- Question: How do we find out which heuristics are best?



Why Simulation?

- Many "HPC" theoretical results
 - Optimal alg., approx. alg., \mathcal{NP} -completeness
- Theory is key to understanding problems better
- Theoretical results also provide motivation and inspiration for developing non-guaranteed heuristics
 - e.g., when there is no optimal algorithm
 - e.g., non-guaranteed algorithms may achieve much better average perfor mance than approximation algorithms
- Sadly, we rarely have theorems to compare heuristics
- Typically, given 20 (reasonable) heuristics and 10,000 random problem instance s, each heuristic could be best on some of the instances
- Question: How do we find out which heuristics are best?



Unrealistic assumptions

- In many cases the "view of the world" of the algorithm is not the real world
 - Example: a heuristic is designed for homogeneous hosts in a cluster, but in practice hosts are a bit heterogeneous
 - Example: network latencies would make the problem NP-complete, an optimal algorithm is known when there are no network latencies, and it can be applied in real-world networks in which there are latencies
 - Example: network topologies are so complex that designing algorithms that truly exploit them is too challenging. Furthermore, in practice one often doesn't know the network topology!
- Question: How do we find out how algorithms behave in the real world?



Unrealistic assumptions

- In many cases the "view of the world" of the algorithm is not the real world
 - Example: a heuristic is designed for homogeneous hosts in a cluster, but in practice hosts are a bit heterogeneous
 - Example: network latencies would make the problem NP-complete, an optimal algorithm is known when there are no network latencies, and it can be applied in real-world networks in which there are latencies
 - Example: network topologies are so complex that designing algorithms that truly exploit them is too challenging. Furthermore, in practice one often doesn't know the network topology!
- Question: How do we find out how algorithms behave in the real world?



Experiment!

- We can't compare heuristics in theory, and in general we don't really know how they would behave in the real world
- So we have to run empirical experiments
 - Create a set of problem instances in the real world
 - Run the heuristics for these instances
 - Do some statistics and try to obtain useful conclusions
- Unfortunately, running experiments is not easy, especially at large scale...



Experiment!

- We can't compare heuristics in theory, and in general we don't really know how they would behave in the real world
- So we have to run empirical experiments
 - Create a set of problem instances in the real world
 - Run the heuristics for these instances
 - Do some statistics and try to obtain useful conclusions
- Unfortunately, running experiments is not easy, especially at large scale...



Experiment!

- We can't compare heuristics in theory, and in general we don't really know how they would behave in the real world
- So we have to run empirical experiments
 - Create a set of problem instances in the real world
 - Run the heuristics for these instances
 - Do some statistics and try to obtain useful conclusions
- Unfortunately, running experiments is not easy, especially at large scale...

The trouble with experiments (I)

■ Experiments are labor-intensive

- To run an experiment on a real-world platform, we need a full-fledge implementation of the studied application/system
- We do not always have such an implementation!
 - e.g., we want to run simulations to decide how to implement the application/system!
 - e.g., we want to answer research questions without necessarily having a particular application at hand with all the required input datasets, etc.
- Developing a full implementation "just" to study scheduling algorithms it not necessarily something one wants to do
 - Or at least not until much later



The trouble with experiments (I)

- Experiments are labor-intensive
- To run an experiment on a real-world platform, we need a full-fledge implementation of the studied application/system
- We do not always have such an implementation!
 - e.g., we want to run simulations to decide how to implement the application/system!
 - e.g., we want to answer research questions without necessarily having a particular application at hand with all the required input datasets, etc.
- Developing a full implementation "just" to study scheduling algorithms it not necessarily something one wants to do
 - Or at least not until much later



The trouble with experiments (II)

- Experiments can be costly in time, \$, and Watts
- As the target platform increases in scale (e.g., large number of hosts) and/or the application increases in scale (e.g., larger data, larger amounts of computations), so does the time for each experiment
- Longer experiments imply larger cost and larger power consumption
- To make matters worse, we typically need large numbers of experiments to achieve reasonable statistical significance

The trouble with experiments (II)

- Experiments can be costly in time, \$, and Watts
- As the target platform increases in scale (e.g., large number of hosts) and/or the application increases in scale (e.g., larger data, larger amounts of computations), so does the time for each experiment
- Longer experiments imply larger cost and larger power consumption
- To make matters worse, we typically need large numbers of experiments to achieve reasonable statistical significance

The trouble with experiments (III)

Experiments are limited in scope

- The set of experimental scenarios is constrained by the platform configurations at hand, and exploring a wide range of configurations may not be possible
- In fact, production platforms may not be available at all, limiting experiments to limited testbeds
- It it difficult to explore hypothetical "what if?" scenarios
 - "What if the network was twice as congested?", "What if network paths were twice as long?", "What if all hosts were 16-core instead of 8-core?"
 - Enabled to some extent by emulation/virtualization techniques, but not possible on all platforms



The trouble with experiments (III)

- Experiments are limited in scope
- The set of experimental scenarios is constrained by the platform configurations at hand, and exploring a wide range of configurations may not be possible
- In fact, production platforms may not be available at all, limiting experiments to limited testbeds
- It it difficult to explore hypothetical "what if?" scenarios
 - "What if the network was twice as congested?", "What if network paths were twice as long?", "What if all hosts were 16-core instead of 8-core?"
 - Enabled to some extent by emulation/virtualization techniques, but not possible on all platforms



The trouble with experiments (IV)

Experiments are not always perfectly repeatable

- As soon as the target platform becomes large (e.g., several large clusters interconnected over wide-area networks) it is typically subject to external and/or unpredictable load conditions
 - A shared network, perhaps even shared hosts
 - Unscheduled downtimes or performance bugs
 - Changing software configurations
- As a result, the same experiment may lead to different results on different days
- The simple approach that consists in running experiments "back-to-back" is not satisfying if each experiment takes a long time



The trouble with experiments (IV)

- Experiments are not always perfectly repeatable
- As soon as the target platform becomes large (e.g., several large clusters interconnected over wide-area networks) it is typically subject to external and/or unpredictable load conditions
 - A shared network, perhaps even shared hosts
 - Unscheduled downtimes or performance bugs
 - Changing software configurations
- As a result, the same experiment may lead to different results on different days
- The simple approach that consists in running experiments "back-to-back" is not satisfying if each experiment takes a long time



Simulation to the rescue

- Given all these difficulties with real-world experiments, a popular approach is simulation
- Simulation: the use of a computer program that implements mathematical and algorithmic models of the behavior of an application running on a platform
- The input to the simulation:
 - A specification of the platform's characteristics
 - A specification of the application's characteristics
- The output of the simulation:
 - A time-stamped list of relevant events throughout (simulated) time
 - From this list can be gathered statistics, visualizations



Simulation to the rescue

- Given all these difficulties with real-world experiments, a popular approach is simulation
- Simulation: the use of a computer program that implements mathematical and algorithmic models of the behavior of an application running on a platform
- The input to the simulation:
 - A specification of the platform's characteristics
 - A specification of the application's characteristics
- The output of the simulation:
 - A time-stamped list of relevant events throughout (simulated) time
 - From this list can be gathered statistics, visualizations



Simulation to the rescue

- Given all these difficulties with real-world experiments, a popular approach is simulation
- Simulation: the use of a computer program that implements mathematical and algorithmic models of the behavior of an application running on a platform
- The input to the simulation:
 - A specification of the platform's characteristics
 - A specification of the application's characteristics
- The output of the simulation:
 - A time-stamped list of relevant events throughout (simulated) time
 - From this list can be gathered statistics, visualizations



The trouble with simulation

- Every simulation introduces a bias, or error
 - The simulation is instantiated with unrealistic parameters
 - The simulation models' implementation is buggy
 - The simulation models are implemented correctly but do not correspond to the real world
 - The simulated application does not correspond to the real-world application because not based on a real implementation
 - The simulation ignores transient real-world behaviors
 - ..
- The larger the simulation error the less useful the simulation
- If the simulation error is not bounded/quantified, then the simulation is even less useful



The trouble with simulation

- Every simulation introduces a bias, or error
 - The simulation is instantiated with unrealistic parameters
 - The simulation models' implementation is buggy
 - The simulation models are implemented correctly but do not correspond to the real world
 - The simulated application does not correspond to the real-world application because not based on a real implementation
 - The simulation ignores transient real-world behaviors
 - ..
- The larger the simulation error the less useful the simulation
- If the simulation error is not bounded/quantified, then the simulation is even less useful



Reducing simulation error (I)

- A (widely accepted) way to reduce simulation error is to have the simulation be highly detailed
- More details can be achieved by making analytical models more complex (i.e., more parameters)
- Example:
 - compute_time = #inst/inst_per_second
 - compute_time = [#inst_{load_store} × (CPI_{load_store} + cache_miss_rate × cache_penalty) + #inst_{other} × CPI_{other}]/clock_rate
 - Accounting for instruction-level parallelism...

Reducing simulation error (I)

- A (widely accepted) way to reduce simulation error is to have the simulation be highly detailed
- More details can be achieved by making analytical models more complex (i.e., more parameters)
- Example:
 - compute_time = #inst/inst_per_second
 - compute_time =
 [#inst_{load_store} × (CPI_{load_store} + cache_miss_rate ×
 cache_penalty) + #inst_{other} × CPI_{other}]/clock_rate
 - Accounting for instruction-level parallelism...

Reducing simulation error (I)

- A (widely accepted) way to reduce simulation error is to have the simulation be highly detailed
- More details can be achieved by making analytical models more complex (i.e., more parameters)
- Example:
 - compute_time = #inst/inst_per_second
 - compute_time =
 [#inst_{load_store} × (CPI_{load_store} + cache_miss_rate ×
 cache_penalty) + #inst_{other} × CPI_{other}]/clock_rate
 - Accounting for instruction-level parallelism...

Reducing simulation error (II)

- More details can be achieved by replacing analytical simulation models by complex programmatic models
- Example:
 - Don't use a formula to estimate the compute time
 - Instead write a program that generates control signals in a hypothetical micro-architecture, simulating what happens at each clock cycles (register set, ALUs, caches, etc.)
 - The input is the actual list of instructions and operands
 - The output is a time-stamped trace of instruction completions from which one can compute statistics
 - So-called *cycle-accurate simulation*



Reducing simulation error (II)

- More details can be achieved by replacing analytical simulation models by complex programmatic models
- Example:
 - Don't use a formula to estimate the compute time
 - Instead write a program that generates control signals in a hypothetical micro-architecture, simulating what happens at each clock cycles (register set, ALUs, caches, etc.)
 - The input is the actual list of instructions and operands
 - The output is a time-stamped trace of instruction completions from which one can compute statistics
 - So-called *cycle-accurate simulation*



The trouble with complex simulation

- The notion that "more complex = better" is not without its problems [Gibson et al., ASPLOS'00]
 - And more complex models may be poorly instantiated
- But making models more complex/sophisticated/complete is still the most common approach
- One problem is: more complex is slower
- The ratio of simulation time to simulated time grows and can become very large (e.g., > 10, > 100)
 - Hardware is fast, (simulation) software is slow
- If simulated scenarios are large/long, then simulation time is prohibitively long
 - Remember that we may need to run thousands of simulations to compare scheduling heuristics



The trouble with complex simulation

- The notion that "more complex = better" is not without its problems [Gibson et al., ASPLOS'00]
 - And more complex models may be poorly instantiated
- But making models more complex/sophisticated/complete is still the most common approach
- One problem is: more complex is slower
- The ratio of simulation time to simulated time grows and can become very large (e.g., > 10, > 100)
 - Hardware is fast, (simulation) software is slow
- If simulated scenarios are large/long, then simulation time is prohibitively long
 - Remember that we may need to run thousands of simulations to compare scheduling heuristics



- We want simulation to be fast and accurate
- It would seem that simulation can be either
 - slow and accurate, or
 - fast and erroneous
- The goal is to push the limits:
 - Improve the speed of "complex" simulation models
 - Improve the accuracy of "simple" simulation models

- We want simulation to be fast and accurate
- It would seem that simulation can be either
 - slow and accurate, or
 - fast and erroneous
- The goal is to push the limits:
 - Improve the speed of "complex" simulation models
 - Improve the accuracy of "simple" simulation models

- We want simulation to be fast and accurate
- It would seem that simulation can be either
 - slow and accurate, or
 - fast and erroneous
- The goal is to push the limits:
 - Improve the speed of "complex" simulation models
 - Improve the accuracy of "simple" simulation models

- We want simulation to be fast and accurate
- It would seem that simulation can be either
 - slow and accurate, or
 - fast and erroneous
- The goal is to push the limits:
 - Improve the speed of "complex" simulation models
 - Improve the accuracy of "simple" simulation models

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

Parallel/Distributed Computing Simulation

- Many results in parallel/distributed computing research are obtained in simulation
- Simulation has been used for decades in various areas of computer science
 - Network protocol design, microprocessor design
- By comparison, current practice in parallel/distributed computing is in its infancy
- Part of the problem is the lack of a standard tool, likely due to research being partitioned in different sub-areas
 - cluster computing, grid computing, volunteer computing, peer-to-peer computing, cloud computing, ...
- Part of the problem is the lack of open methodology...

Lack of open methodology

- The key to open science is that results can be reproduced
- Unfortunately, open science is not yet possible in parallel/distributed computing simulations
- Example: "Towards Yet Another Peer-to-Peer Simulator", Naicken et al., *Proc. of HET-NETs*, 2006
 - Surveyed 141 papers that use simulation
 - 30% use a custom simulator
 - 50% don't say which simulator is used (!!)
 - Only few authors provide downloadable software artifacts / datasets
- Researchers don't trust simulators developed by others?
- They feel they can develop the best simulator themselves?
- Ironic consequence: published simulation results are often not reproducible



Lack of open methodology

- The key to open science is that results can be reproduced
- Unfortunately, open science is not yet possible in parallel/distributed computing simulations
- Example: "Towards Yet Another Peer-to-Peer Simulator", Naicken et al., Proc. of HET-NETs, 2006
 - Surveyed 141 papers that use simulation
 - 30% use a custom simulator
 - 50% don't say which simulator is used (!!)
 - Only few authors provide downloadable software artifacts / datasets
- Researchers don't trust simulators developed by others?
- They feel they can develop the best simulator themselves?
- Ironic consequence: published simulation results are often not reproducible

Lack of open methodology

- The key to open science is that results can be reproduced
- Unfortunately, open science is not yet possible in parallel/distributed computing simulations
- Example: "Towards Yet Another Peer-to-Peer Simulator", Naicken et al., Proc. of HET-NETs, 2006
 - Surveyed 141 papers that use simulation
 - 30% use a custom simulator
 - 50% don't say which simulator is used (!!)
 - Only few authors provide downloadable software artifacts / datasets
- Researchers don't trust simulators developed by others?
- They feel they can develop the best simulator themselves?
- Ironic consequence: published simulation results are often not reproducible



And yet....

- There are parallel/distributed computing simulators that are intended to be used by others!
- These simulators all attempt to do the same 3 things:
 - Simulate CPUs
 - Simulate networks
 - Simulate storage
- The two main concerns are:
 - Simulation accuracy
 - Simulation speed / scalability
 - Ability to simulate large/long-running applications/platforms fast and without too much RAM
- Simulators make choices to trade off one for the other
- Let's look at typical such choices...

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

Simulating computation (I)

- There are two main approaches for simulating computations: microscopic or macroscopic
- Microscopic approach: cycle-accurate simulation
 - Simulate micro-architecture components at the clock cycle resolution based on instructions from real (compiled) code or for synthetic instruction mixes
 - Typically used by computer architecture researchers
 - Arguably very accurate (bus contention, cache effects, instruction-level parallelism, multi-core, GPUs, ...)
 - © Very slow (ratio of simulation/simulated time > 100)
 - Arbitrary unscalable as the volume of computation increases, which is a problem for simulating long-running applications (e.g., grid computing, cluster computing)

Simulating computation (II)

- Macroscopic approach: (scaled) delays
 - Defined for each compute resource a *computation speed*
 - Define a computation as a computation volume
 - The simulated time is computed as the ratio of the two, plus an optional random component
 - Reasonably accurate for compute-bound applications
 - Compute times can also be sampled from real application or benchmark executions on a reference computer architecture, and then scaled
 - \odot Scalable: simulation of a computation in O(1) space/time
 - Can be wildly inaccurate for memory-bound applications

Simulating communication

- Each link: latency and bandwidth
- Protocol-accurate packet-level simulation
 - Used by network protocol researchers
 - Accurate
 - slow and not scalable
- Non-protocol-accurate store-and-forward of packets
 - An attempt to be more scalable than the above, but can be made arbitrarily inaccurate (packet size?)
- Ad-hoc fluid models
 - Scalable
 - On the protocol of the prot
- TCP fluid models
 - Scalable and accurate within limits

Simulating storage

- Microscopic approach: detailed discrete-event simulation (e.g., DiskSim)
 - Accurate
 - Just like cycle-accurate, and packet-level: unscalable
 - Arbitrarily high simulated/simulation ratio for large data
- Most simulators provide very little
 - Notion of storage capacity and stored data
 - Each transfer has a latency and bandwidth, with optional random components
 - Fails to capture caching effects, file system effects, ...

Specifying the simulated application (I)

■ There are three main approaches: Finite automata, Event traces, Concurrent Sequential Processes (CSP)

Finite automata:

- Each simulated process is described as an automaton or a Markov chain
- Each state is an action that lasts for some number of (simulated) seconds
 - Example: compute for 10s, then with probability 0.5 communicate for 1s, then with probability 1.0 do nothing for 30s, ...
- Very scalable since each process is described with only a few bytes and fast algorithms can do state transitions
- Limited expressive power, no/little application logic

Specifying the simulated application (II)

Event traces:

- Each simulated process is described as a sequence of compute, communicate, store events
 - At t = 0 compute for 10s, at time t = 10 send a message for 2s, at t = 12 receive a message for 20s, ...
- Events obtained from real-world executions and "replayed", while scaling delays
- Fast since event replay is algorithmically simple
- Not always scalable (traces can be large, obtaining large traces on dedicated platforms can be hard)
- Difficult to extrapolate traces to simulate more/fewer processes

Specifying the simulated application (III)

- Concurrent Sequential Processes (CSP):
 - Application implemented as fragments of arbitrary code that call simulation API functions
 - sim_compute, sim_send, sim_recv, ...
 - Maximum expressive power, arguably
 - Not scalable
 - Number of threads is limited (e.g., "only" a few thousands Java threads)
 - Context switching and synchronization overhead can be high
 - The user can implement expensive logic/computation, which may be useful but also unscalable

State-of-the-art simulators

Simulator	Community of Origin					Application			Network						CPU		Disk		
	high perf. comp.	grid comp.	cloud comp.	volunteer comp.	peer-to-peer comp.	execution trace	abstract spec.	programmatic spec.	latency	bandwidth	store-and-forward	ad-hoc flow	TCP fluid model	packet-level	scaled real delay	scaled delay	capacity	seek + transfer	block-level
PSINS	√					√			√	√		✓			√	√			
LogGOPSim	✓					✓			V	✓						√			
BigSim	✓					✓			V	√				✓	V	√			
MPI-SIM	✓							✓	√	✓					✓	√			
OptorSim		✓						✓	V	✓		✓				√	✓		
GridSim		✓						✓	√	\	✓	✓				√	✓	✓	
GroudSim		✓	✓					✓	√	~		√				~			
CloudSim			√					√	V	✓	✓					√	√	✓	
iCanCloud			✓					✓	✓	✓				✓		√	✓	✓	✓
SimBA				√			√	✓	√							√			
EmBOINC				✓			✓	✓	✓							√			
SimBOINC				✓				✓	√	~			✓			V			
PeerSim					√		√		V										
OverSim					✓			✓	✓	✓				✓					
SIMGRID		√				√	√	√	V	√			√	√	V	√	V	√	

SIMGRID: a counter-intuitive design approach

- Accepted wisdom: to be both accurate and fast/scalable, a simulator must be highly specialized to a target domain
- Typical rationale: to achieve scalability one must "cut corners" and reduce accuracy in ways that are hopefully ok for the target domain
 - Example: when simulating a p2p application, no need to simulate network contention, or compute times
 - Example: when simulating a cluster computing application, no need to simulate external load on the system
- Instead, with SIMGRID we target multiple domains:
 - Grid, cloud, cluster/HPC, volunteer, peer-to-peer
- We claim/demonstrate that we can be accurate/scalable across domains!

SIMGRID: a counter-intuitive design approach

- Accepted wisdom: to be both accurate and fast/scalable, a simulator must be highly specialized to a target domain
- Typical rationale: to achieve scalability one must "cut corners" and reduce accuracy in ways that are hopefully ok for the target domain
 - Example: when simulating a p2p application, no need to simulate network contention, or compute times
 - Example: when simulating a cluster computing application, no need to simulate external load on the system
- Instead, with SIMGRID we target multiple domains:
 - Grid, cloud, cluster/HPC, volunteer, peer-to-peer
- We claim/demonstrate that we can be accurate/scalable across domains!

Outline

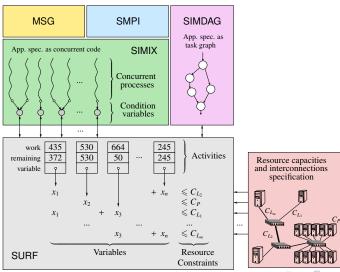
- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

SIMGRID: history

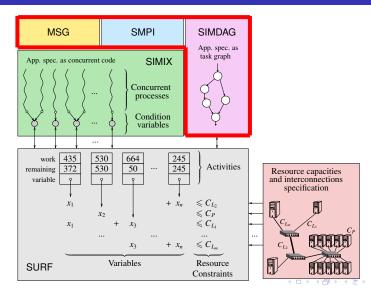
- A 14-year old open source project
- First release: 1999
 - A tool to prototype/evaluate scheduling heuristics, with naïve resource models
 - Essentially: a way to not have to draw Gantt charts by hand
- Second release: 2000
 - Addition of TCP fluid network models
 - Addition of an API to describe simulated app as CSPs
- Third release: 2005
 - Stability, documentation, packaging, portability, ...
- Release v.3.3: 2009
 - Complete rewrite of the simulation core for better scalability
 - Possible to describe transient resource behavior via traces
 - Addition of an "Operating System"-like layer
 - Two new APIs



Software stack



APIs



The SIMDAG API

- Application is described as a task graph
 - SD_TASK_CREATE(), SD_TASK_DEPENDENCY_ADD(), ...
 - SD_TASK_SCHEDULE() (on a host)
 - All types of API functions to get/set task properties/parameters
- One call to SD_SIMULATE() launches the simulation:
 - While there are ready tasks, run them
 - Computation + communication
 - Resolve dependencies
 - Repeat
- A very simple API designed for users who don't need the full power of the CSP abstraction
 - Looks a lot like SIMGRID v1.0

The SMPI API

- Designed to simulate Message Passing Interface (MPI) applications
 - Standard way to implement communication in parallel applications
- The (almost unmodified) application is compiled so that MPI processes run as threads
- MPI calls are intercepted and passed to SIMGRID's simulation core
- "Tricks" are used to allow simulation on a single computer
 - CPU burst durations are samples a few times and "replayed" for free
 - Arrays are shared among threads (wrong data-dependent application behavior but small memory footprint)

The MSG API

- This is the most commonly used API: basic CSP abstraction
- It has bindings in C, Java, Lua, Ruby
- Let's go through a full (but simple) *master-worker* example
 - The master process has tasks to send to workers
 - Each worker "processes" the tasks until it receive a termination signal
- Let's look at:
 - Master code
 - Worker code
 - Main function
 - XML platform description file
 - XML application deployment description file

Master

int master(int argc, char *argv[]) {

```
int number of tasks = atoi(argv[1]); double task comp size = atoi(argv[2]);
double task comm size = atof(argv[3]); int workers count = atoi(argv[4]);
char mailbox[80]; int i;
char buff [64];
msg_task_t task;
/* Dispatching tasks (dumb round-robin algorithm) */
for (i = 0; i < number of tasks; i++) {
  sprintf(buff, "Task %d", i);
  task = MSG task create (buff, task comp size, task comm size, NULL);
  sprintf(mailbox, "worker-%d", i % workers count);
  print("Sending task %s to mailbox %s\n", buff, mailbox);
  MSG task send (task, mailbox);
/* Send finalization message to workers */
for (i = 0; i < workers count; i++)</pre>
  sprintf(mailbox, "worker-%ld", i % workers count);
  MSG task send (MSG task create ("finalize", 0, 0, 0), mailbox);
return 0:
```

Worker

int worker(int argc, char *argv[]) {

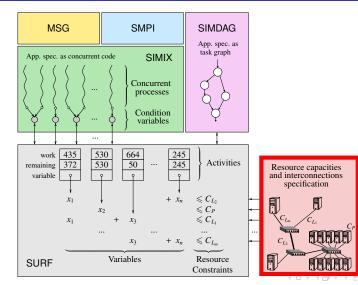
```
msg task t task; int errcode; int id = atoi(argv[1]);
char mailbox[80];
sprintf (mailbox, "worker-%d", id);
while(1) {
  /* Receive a task */
  errcode = MSG task receive (&task, mailbox);
  if (errcode != MSG OK) {
    print("Error"); return errcode;
  if (!strcmp(MSG task get name(task), "finalize")) {
    MSG task destroy (task);
    break:
  print("Processing %s", MSG task get name(task));
  MSG task execute (task):
  print("Task %s done", MSG task get name(task));
  MSG task destroy (task);
print("Worker done!");
return 0;
```

Main program

int main(int argc, char *argv[]) {

```
char *platform file = "my platform.xml";
char *deployment file = "mv deployment.xml";
MSG init (&argc, argv);
/* Declare all existing processes, binding names to functions */
MSG function register ("master", &master);
MSG function register ("worker", &worker);
/* Load a platform description */
MSG create environment (platform file);
/* Load an application deployment description */
MSG launch application (deployment file);
/* Launch the simulation (until its terminates) */
MSG main();
print("Simulated execution time %g seconds", MSG get clock());
```

Platform and app deployment description



Platform description file (XML)

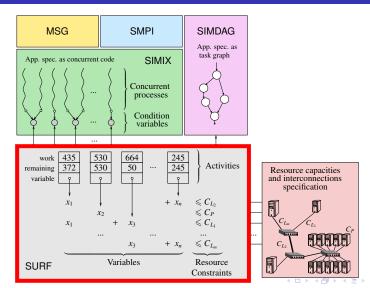
my_platform.xml

```
<?xml version="1 0"?>
<!DOCTYPE platform SYSTEM "http://simgrid.gforge.inria.fr/simgrid.dtd">
  <platform version="3">
    <AS id="mynetwork" routing="Full">
      <host id="host1" power="1E6" />
      <host id="host2" power="1E8" />
      <host id="host3" power="1E6" />
      <host id="host4" power="1E9" />
      <link id="link1" bandwidth="1E4" latency="1E-3" />
      link id="link2" bandwidth="1E5" latency="1E-2" />
      link id="link3" bandwidth="1E6" latency="1E-2" />
      <link id="link4" bandwidth="1E6" latency="1E-1" />
      <route src="host1" dst="host2"> <link id="link1"/> <link id="link2"/> </route>
      <route src="host1" dst="host3"> <link id="link1"/> <link id="link3"/> </route>
      <route src="host1" dst="host4"> <link id="link1"/> <link id="link4"/> </route>
    </AS>
  </platform>
```

Application deployment description file (XML)

my_deployment.xml

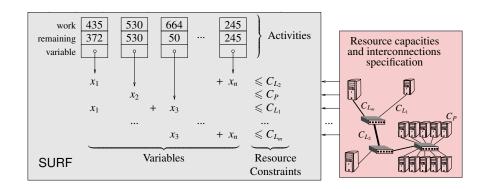
Simulation core



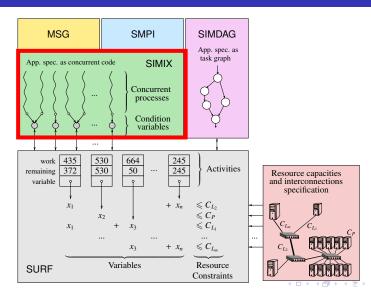
Simulation core

- The SURF component implements all simulation models
- All application activities are called actions
- SURF keeps track of all actions:
 - Work to do
 - Work that remains to be done
 - Link to a set of variables
- All action variables occur in constraints
 - Capture the fact that actions use one of more resources
- SURF solves the a (modified) Max-min constrained optimization problem betwen each simulation event
 - See previous seminar for more details
- Let's explain the example in the figure...

Simulation core



The SIMIX simcall interface



The SIMIX simcall interface

- SIMIX: an "OS kernel" on top of the simulation core
- Each simulated processes is a "thread" (more on this later)
- These threads run in mutual exclusion, round-robin, as controlled by SIMIX
- Each time a thread places an API call, translated to a simcall (a simulated syscall), it blocks on a condition variable in SIMIX
- When all threads are blocked in this way, SIMIX tells the simulation core computes the simulation models
- Threads are then unblocked and proceed until they all enter the SIMIX "kernel" again
- Total separation: application / synchronization / models

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

Scalability

- We have spent many years trying to increase scalability
- The first step was fast analytical resource modeling
 - Solving a weighted Max-min problem as opposed to packet-level, cycle-accurate simulation
 - Implementing the solver with cache-efficient data structures
- Scalability issues in SIMGRID don't come from the models!
- Four limits to scalability:
 - X Running the simulation models too often
 - X Too large platform descriptions
 - **X** Too many simulated processes
 - X Simulation limited to a single core

Scalability

- We have spent many years trying to increase scalability
- The first step was fast analytical resource modeling
 - Solving a weighted Max-min problem as opposed to packet-level, cycle-accurate simulation
 - Implementing the solver with cache-efficient data structures
- Scalability issues in SIMGRID don't come from the models!
- Four limits to scalability:
 - **X** Running the simulation models too often
 - **X** Too large platform descriptions
 - X Too many simulated processes
 - X Simulation limited to a single core

Running models too often

- SIMGRID was originally intended for simulating tightly-coupled parallel apps on hierarchical networks
 - e.g., a grid platforms with 3 clusters on a fast wide-area network for running parallel scientific applications
- In this setting, every simulated action can have an impact on every other simulated action
 - A data transfer completion frees up some bandwidth usable by many other transfers
 - A computation completion can lead to a message that will unblock many other computations
- As a result, SIMGRID was implemented in the typical loop:
 - Run all simulation models
 - Determining the next event (e.g., action completions)
 - Update all actions remaining work amounts
 - Advance the simulated time and repeat



Running models too often

- As SIMGRID gained popularity, we and users tried to apply it to different domains
- One such domain: volunteer computing
 - Donated compute cycles and disk space at the edge of the network to contribute to public-interest projects
 - e.g., SETI@Home, AIDS@Home, BOINC, etc.
- In this setting, many actions are independent
 - There is little resource contention among participating hosts
 - Computations are independent and long-running
- Yet there are many events (thousands of simulated hosts)
- Essentially, SIMGRID keeps decreasing the remaining work amounts of all actions by ε over and over
- The result: sloooooow simulations at large-scale



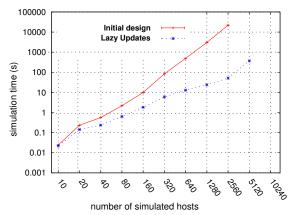
Lazy action updates

- Modified "Lazy Updates" simulation loop:
 - All actions are stored in a heap, sorted by their current completion dates
 - When a resource state is modified, we remove relevant actions (those that use the resource) from the heap, we update their remaining work amounts and completion dates, and we re-insert them into the heap
 - Removing/Inserting from/to a heap: $O(\log n)$
 - Finding the next action that completes: O(1)
- Not a revolutionary idea of course
 - Large simulation literature on efficient *future event sets*
 - But not seen in parallel/distributed computing simulators
- If the application is tightly coupled, then lazy updates are slower because all actions are removed/inserted
- Lazy updates enabled by default but optional



Lazy updates in action

 Lazy updates for the motivating volunteer computing scenario [Heien et al., 2008]



SIMGRID better than specialized simulator?

- Lazy updates are really effective
 - Example: from 3h to 1min for a simulation with 2,560 hosts on a 2.2GHz processor
- And in fact, SIMGRID, even though it implements more sophisticated (network) models is ~25 times faster than the SimBA volunteer computing simulator
- SimBA was optimized for scalability in a different way
 - It uses finite automata to describe simulated processes
- And yet, a more versatile simulator can "out-scale" it thanks to careful design

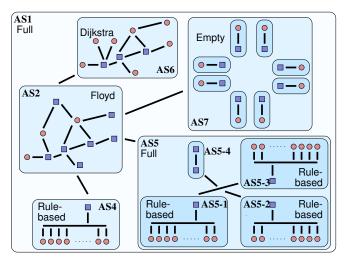
Four limits to SIMGRID's scalability

- √ Running the simulation models too often
- **X** Too large platform descriptions
- **X** Too many simulated processes
- X Simulation limited to a single core

Large platform descriptions

- Users who used SIMGRID for truly large-scale platform simulations often found themselves stuck
 - Long XML parse time
 - Out-of-memory errors
 - Long time to compute network routes, especially because we need $N \times (N-1)$ routes for N hosts!
- To enable large-scale simulation we must have hierarchical platform descriptions
- A platform is an Autonomous System (AS), that can contain interconnected ASes
- Each AS has its own routing scheme
 - Full routing tables, Dijkstra, Floyd, no routing, routing based on rules encoded as regular expressions

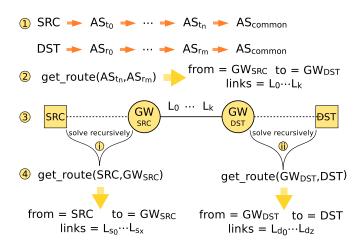
Platform description example



Recursive route computation

- Each AS declared gateways to other ASes, and that are used to compute routes
- To determine a route between two hosts:
 - Search for the common ancestor AS Search for the route between the two relevant AS children
 - Repeat until the full route is determined
- Let's see this on a figure...

Recursive route computation



Generating platform descriptions

- The SIMGRID user can either generate an XML file, or used SIMGRID's platform generation API
- The overhead of (recursively) computing the route is negligible in our implementation
- The memory footprint of the platform description is small
- XML parsing is fast
- Example:
 - The Grid'5000 testbed (10 sites, 40 clusters, 1,500 nodes)
 - Described with 22KiB of XML, parsed in < 1s</p>
 - Previous SIMGRID versions: 520MiB, parsed in 20min

Four limits to SIMGRID's scalability

- √ Running the simulation models too often
- √ Too large platform descriptions
- X Too many simulated processes
- X Simulation limited to a single core

Too many threads

- SIMGRID allows users to described simulated apps as sets of CSPs
- Great for flexibility and expressivity
- Not scalable if implemented as processes/threads:
 - Thread creation/management overhead in the kernel
 - Thread memory footprint in the kernel
 - Thread synchronization overhead (locks + condvar)
- But in a SIMGRID simulation threads run in mutual exclusion and in a round-robin fashion
- Therefore, we don't need the full power/flexibility of kernel threads since we do our own scheduling and our own synchronization

Scalable CSPs (I)

- As opposed to having threads each with a bunch of locks and condition variables we take a different approach
- A single "core context":
 - Since simulation models are fast, a single thread does all model computations (i.e., it run the SURF code)
 - All simulated processes place SIMIX simcalls, and all these simcalls are resolved by the core context: no shared state among threads
 - Two simulated processes waiting on each other don't really wait on each other
 - i.e., no multi-step process-to-process interactions
 - Instead, they place wait/notify-like simcalls to the core context

Scalable CSPs (I)

■ Lightweight "continuations":

- Since we don't need full threads we can use cooperative, light-weight, non-preemptive threads
 - Known as continuations
 - No actual context-switching by the kernel
- Windows: fibers
- Linux, Mac OSX: ucontexts
- We actually re-implemented them in assembly to avoid a costly system call that is not needed for our purpose

How scalable is it?

- With all three scalability improvements so far, we can now compare SIMGRID to "competitors"
- Case study #1: Grid computing
 - Master-worker scenario
 - Comparison to GridSim (implemented in Java)
- Case study #2: Peer-to-peer computing
 - The Chord protocol [Stoica et al., 2003]
 - Comparison to PeerSim and OverSim

How scalable is it?

- With all three scalability improvements so far, we can now compare SIMGRID to "competitors"
- Case study #1: Grid computing
 - Master-worker scenario
 - Comparison to GridSim (implemented in Java)
- Case study #2: Peer-to-peer computing
 - The Chord protocol [Stoica et al., 2003]
 - Comparison to PeerSim and OverSim

How scalable is it?

- With all three scalability improvements so far, we can now compare SIMGRID to "competitors"
- Case study #1: Grid computing
 - Master-worker scenario
 - Comparison to GridSim (implemented in Java)
- Case study #2: Peer-to-peer computing
 - The Chord protocol [Stoica et al., 2003]
 - Comparison to PeerSim and OverSim

- One master, N workers, P tasks, round-robin scheduling
- Simulation on a 2.4GHz core and 8GiB of RAM
- GridSim:
 - No network topology simulated (simply latency+bandwidth communication costs)
- SimGrid:
 - Grid'5000 topology simulated (with TCP flow-level modeling, etc.)

Polynomial fits based on measured values

	Simulation Time (s)	Peak Memory Footprint (byte)
GridSim	$5.599 \times 10^{-2}P + 1.405 \times 10^{-8}N^2$	$2.457 \times 10^6 + 226.6P + 3.11N$
SIMGRID	$1.021 \times 10^{-4}P + 2.588 \times 10^{-5}N$	5188 + 79.9P

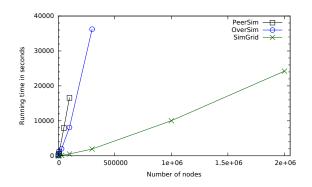
Example: N = 2,000, P = 500,000

GridSim: 4.4GiB of RAM, > 1hour

■ SIMGRID: 165MiB of RAM, < 14s

And SIMGRID uses more sophisticated models!

- Implementation of the Chord protocol for N hosts
- Simulations on a 1.7Ghz core with 48GiB of RAM
- SIMGRID
 - TCP flow-level modeling on a full topology
- OverSim [Baumgart et al., 2007]
 - Communication delays based on Euclidian distance between peers
 - Implemented in C++
- PeerSim [Montresor et al., 2009]
 - Constant communication delays
 - Implemented in Java



■ PeerSim: 100,000 peers in 4h36min

OverSim: 200,000 peers in 10h

■ SIMGRID: 2,000,000 peers in 32min

Scalability results

- We have shown SIMGRID to be faster than specialized simulators, even when it uses more sophisticated network simulation models!
- Some of these simulators were designed specifically for scalability (especially p2p simulators)
 - Of course, they may suffer from implementation inefficiencies, while we have spent hours trying to optimize our implementation
- Nevertheless, we claim that it is not necessary to be specialized to be scalable, at least for parallel/distributed computing simulations
- Can we go further?

Scalability results

- We have shown SIMGRID to be faster than specialized simulators, even when it uses more sophisticated network simulation models!
- Some of these simulators were designed specifically for scalability (especially p2p simulators)
 - Of course, they may suffer from implementation inefficiencies, while we have spent hours trying to optimize our implementation
- Nevertheless, we claim that it is not necessary to be specialized to be scalable, at least for parallel/distributed computing simulations
- Can we go further?

Four limits to SIMGRID's scalability

- √ Running the simulation models too often
- √ Too large platform descriptions
- √ Too many simulated processes
- **X** Simulation limited to a single core

Parallelizing SIMGRID

- Because of all the optimizations we've talked about, often most of the compute time is spent in user code!
 - What simulated processes do outside of SIMGRID
- There is thus no need to parallelize SIMGRID's internals
 - Which would be very difficult anyway since Parallel Discrete
 Event Simulation is difficult
- We are thus able to run concurrent user processes easily on multiple cores
- Experiments for "difficult cases" (e.g., peer-to-peer Chord) show that achieved speedup is minimal (13%) but non-zero
- Experiments for "easy cases" (e.g., simulated processes that do complex logic in between calls to SIMGRID) show that achieved speedup up is large

Parallelizing SIMGRID

- Because of all the optimizations we've talked about, often most of the compute time is spent in user code!
 - What simulated processes do outside of SIMGRID
- There is thus no need to parallelize SIMGRID's internals
 - Which would be very difficult anyway since Parallel Discrete
 Event Simulation is difficult
- We are thus able to run concurrent user processes easily on multiple cores
- Experiments for "difficult cases" (e.g., peer-to-peer Chord) show that achieved speedup is minimal (13%) but non-zero
- Experiments for "easy cases" (e.g., simulated processes that do complex logic in between calls to SIMGRID) show that achieved speedup up is large

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

Network simulation

- An important aspect of the execution of a parallel and distributed application is the time spent doing network activities
 - Sending messages
 - Waiting for messages
- Many proposed scheduling algorithms attempt to carefully schedule both computation and communication
- But the "view of the network" of the algorithms is notoriously different from the reality of deployed networks

Packet-level simulation

- The network community, i.e., researchers concerned with the design of network protocols, routing schemes, etc., typically use packet-level simulation
- A packet-level simulator is a discrete-event simulator
 - packet emission/reception events
 - network protocol events
- The life-cycle of every individual packet is simulated
 - e.g., from the TCP stack down the the IP level
- Popular simulators: NS2, NS3, GTNetS, OMNet++, ...
 - Some use simplified protocol stacks (GTNeTS, NS2)
 - Some can use real TCP implementations (NS3)
- In this seminar: simulation of wired TCP/IP networks

Packet-level simulation: accurate

- Since packet-level simulation simulates every packet and implements protocol stack, it is often accepted as accurate
 - Or at least, published packet-level simulation results are typically trusted
- It is outside the scope of this seminar to discuss to which extent this is true
 - It turns out there not many validation studies
 - And the word "accurate" means different things to different people
 - e.g., matching the spec of how we design the network
 - vs. matching what actually happens in a real-world environment
- We will simply assume that "packet-level = perfectly accurate"

Packet-level simulation: accurate

- Since packet-level simulation simulates every packet and implements protocol stack, it is often accepted as accurate
 - Or at least, published packet-level simulation results are typically trusted
- It is outside the scope of this seminar to discuss to which extent this is true
 - It turns out there not many validation studies
 - And the word "accurate" means different things to different people
 - e.g., matching the spec of how we design the network
 - vs. matching what actually happens in a real-world environment
- We will simply assume that "packet-level = perfectly accurate"

Packet-level simulation: slow

- High simulation/simulated ratio
 - Using GTNetS, which is known for being reasonable fast
 - Simulating 200 communications each transferring 100MiB between two random end-points in a random 200-node topology
 - 125 sec of simulated time, 1,500 sec of simulation time on a 3.2GHz processor
 - The ratio can be made arbitrarily large by increasing the number of communications
- This is because each packet leads to many network events to be simulated in software
- Not usable to run (many) simulations in which there are large numbers of communications that involve large numbers of packets

Packet-level simulation: slow

- High simulation/simulated ratio
 - Using GTNetS, which is known for being reasonable fast
 - Simulating 200 communications each transferring 100MiB between two random end-points in a random 200-node topology
 - 125 sec of simulated time, 1,500 sec of simulation time on a 3.2GHz processor
 - The ratio can be made arbitrarily large by increasing the number of communications
- This is because each packet leads to many network events to be simulated in software
- Not usable to run (many) simulations in which there are large numbers of communications that involve large numbers of packets

An alternative to packet-level simulation

- To make simulation fast, we must avoid simulating individual packets
- Simple idea: simulate flows, abstracting away packets, and determine flow completion times via equations
 - flow: source + network path + destination + #bytes (S)
- The simplest of flow models: T = latency + bandwidth/S
 - latency = end-to-end latency (sec)
 - bandwidth = bottleneck bandwidth (byte/sec)
- There are, unfortunately, several problems with the above
 - Network protocol overhead?
 - Network protocol effects (e.g., congestion window)?
 - How do we find out the bottleneck bandwidth???
- Let's look at some flow-level simulators

An alternative to packet-level simulation

- To make simulation fast, we must avoid simulating individual packets
- Simple idea: simulate flows, abstracting away packets, and determine flow completion times via equations
 - flow: source + network path + destination + #bytes (S)
- The simplest of flow models: T = latency + bandwidth/S
 - latency = end-to-end latency (sec)
 - bandwidth = bottleneck bandwidth (byte/sec)
- There are, unfortunately, several problems with the above
 - Network protocol overhead?
 - Network protocol effects (e.g., congestion window)?
 - How do we find out the bottleneck bandwidth???
- Let's look at some flow-level simulators

Some simulators

- A research area in which non-packet-level network simulation has been used for decades is "Grid computing" (and lately "Cloud computing")
 - Clusters federated via wide-area networks, many processes, large data
- Four well-known, and/or popular, and/or widely used simulators that employ flow-level models
 - GridSim [Buyya et al., CCPE'02] (~ 10 download/day)
 - OptorSim [Bell et al., IJHPCA'03]
 - GroudSim [Ostermann at al., Grid'10]
 - CloudSim [Calheiros et al., SPE'11]
- These are MUCH faster than packet-level simulators
- But are they accurate?

Some simulators

- A research area in which non-packet-level network simulation has been used for decades is "Grid computing" (and lately "Cloud computing")
 - Clusters federated via wide-area networks, many processes, large data
- Four well-known, and/or popular, and/or widely used simulators that employ flow-level models
 - GridSim [Buyya et al., CCPE'02] (~ 10 download/day)
 - OptorSim [Bell et al., IJHPCA'03]
 - GroudSim [Ostermann at al., Grid'10]
 - CloudSim [Calheiros et al., SPE'11]
- These are MUCH faster than packet-level simulators
- But are they accurate?

Some simulators

- A research area in which non-packet-level network simulation has been used for decades is "Grid computing" (and lately "Cloud computing")
 - Clusters federated via wide-area networks, many processes, large data
- Four well-known, and/or popular, and/or widely used simulators that employ flow-level models
 - GridSim [Buyya et al., CCPE'02] (~ 10 download/day)
 - OptorSim [Bell et al., IJHPCA'03]
 - GroudSim [Ostermann at al., Grid'10]
 - CloudSim [Calheiros et al., SPE'11]
- These are MUCH faster than packet-level simulators
- But are they accurate?

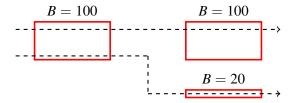
Invalidating experiments

- It turns out these simulators are not accurate
 - In spite of perhaps convincing published results in some cases
- Very easy to come up with invalidating experiments
- Let's exhibit such experiments with the following visuals:

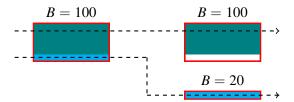


A flow receiving a bandwidth share of a link of bandwidth B

Invalidating OptorSim and GroudSim

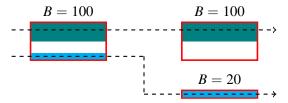


Invalidating OptorSim and GroudSim



What a real network would do (link utilization is maximized)

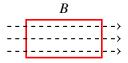
Invalidating OptorSim and GroudSim



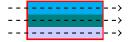
What OptorSim and GroudSim do (links are shared 50% in spite of bottlenecks)

This "weakness" is document by the authors

Invalidating GridSim

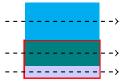


Invalidating GridSim



What a real network would do

Invalidating GridSim



What GridSim does (first flow receives B, second B/2, etc.)

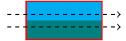
Likely a bug (which we reported)

Invalidating CloudSim

One flow starts ϵ seconds after the other

Invalidating CloudSim

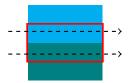
One flow starts ϵ seconds after the other



What a real network would do

Invalidating CloudSim

One flow starts ϵ seconds after the other

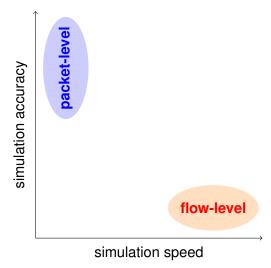


What CloudSim does (each flow receives *B*)

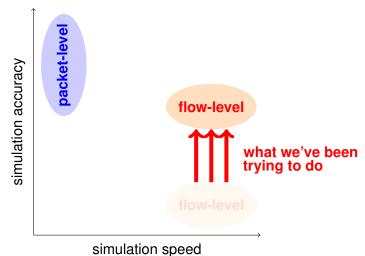
Invalidation is easy???

- We took four well-known simulators
- Suggested by our own experience and by inspecting the code of those simulators, we came up with trivial invalidating experiments
- It is quite shocking how easy it was to perform the invalidation
- One may wonder: how is this possible?
- Answer: validation studies are rare (and invalidation studies are more rare)
 - Published results typically show "good cases"

Is there no hope then?



Is there no hope then?



Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

The two main questions

- Consider a network topology (end-points, routers, links), latency and bandwidth values for the links, and a set of flows each with some start time
- The two questions we need to answer are:
 - 1. What bandwidth is allocated to each flow throughout its execution?
 - 2. What is the completion time of each flow?
- Question 1 can be answered (approximately) if we take a steady-state simulation approach

Steady-state bandwidths

- We assume that each flow *f* present at time *t* run forever, and that no new flow arrives in the system
- \blacksquare Given these assumptions, we compute the constant bandwidth allocated to each flow, B_f
 - In a real network the bandwidths would converge reasonably quickly to pseudo-constant values
- For each flow f, given S_f (remaining amount of data) and B_f , we compute the flow's hypothetical completion time t_f
- Say the next new flow arrives at time t_{new}
- Between time t and time $t' = \min(t_{new}, \min_f t_f)$ the number of flows is constant, and we know their bandwidths!
- We "advance" the simulation to time t', updating $S_f \leftarrow S_f B_f/(t'-t)$

Bandwidth constraints

With the steady-state approach, i.e., no notions of time and flow sequencing, we can formalize the bandwidth computation problem nicely

Problem (Constraints)

- \blacksquare \mathcal{F} a set of flows, \mathcal{L} a set of links
- B_l is the bandwidth of $l \in \mathcal{L}$
- lacksquare ho_f is the bandwidth allocated to $f \in \mathcal{F}$

$$\forall l \in \mathcal{L}, \sum_{f \in \mathcal{F} \text{ going through } l}
ho_f \leq B_l$$

Objective: compute realistic ρ_f 's

Computing realistic bandwidths

- The question of computing realistic bandwidths has been studied in the context of TCP
- Two main approaches have been developed
- 1. Bottom-up:
 - Reason on the microscopic behavior of TCP
 - Infer its macroscopic behavior
- 2. Top-down:
 - Propose a reasonable model of macroscopic behavior
 - Instantiate it based on (in)validation experiments

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

A bottom-up model

- TCP uses a congestion window that bounds the number of in-flight packets for a flow, so as to perform congestion control at the end-points
- The window size is tuned via additive increase and multiplicative decrease
 - Ramp up slowly, back down fast
- Several authors have proposed bottom-up "bandwidth sharing" models by reasoning on window size dynamics
 - Mo et al. [INFOCOM'99], [IEEE TN 2000]
 - Yaïche et al. [IEEE TN 2010]
 - Low et al. [J. ACM 2002], [IEEE TN 2003]

A bottom-up model

- TCP uses a congestion window that bounds the number of in-flight packets for a flow, so as to perform congestion control at the end-points
- The window size is tuned via additive increase and multiplicative decrease
 - Ramp up slowly, back down fast
- Several authors have proposed bottom-up "bandwidth sharing" models by reasoning on window size dynamics
 - Mo et al. [INFOCOM'99], [IEEE TN 2000]
 - Yaïche et al. [IEEE TN 2010]
 - Low et al. [J. ACM 2002], [IEEE TN 2003]

- $w_f(t)$: the window size for flow $f \in \mathcal{F}$ at time t
- d_f : the equilibrium RTT (propagation plus equilibrium queuing delay) of f, which is assumed to be constant
- f's instantaneous data transfer rate in packets per time unit is $\rho_f(t) = w_f(t)/d_f$
- $\blacksquare q_f(t)$: the packet loss probability for flow f at time t
- At time t, f's emitter transmits $\rho_f(t)$ packets per time units
- For these packets it receives $\rho_f(t)(1-q_f(t))$ positive acks and $\rho_f(t)q_f(t)$ negative acks
- Additive increase: the window increased by $1/w_f(t)$
 - lacksquare $w_f(t)$ is increased by 1 for each $w_f(t)$ packets sent
- Multiplicative decrease: the window is halved

■ The net change in window size per time unit is:

$$w_f(t+1) - w_f(t) = \rho_f(t) \frac{1}{w_f(t)} (1 - q_f(t)) - \rho_f(t) \frac{w_f(t)}{2} q_f(t)$$

■ Since $w_f(t) = \rho_f(t) \times d_f$, we obtain

$$\rho_f(t+1) - \rho_f(t) = \frac{1}{d_f^2} (1 - q_f(t)) - \frac{1}{2} \rho_f(t)^2 q_f(t)$$

As an approximation we obtain a differential equation:

$$\frac{\partial \rho_f}{\partial t}(t) = \frac{1}{d_f^2}(1 - q_f(t)) - \frac{1}{2}\rho_f(t)^2 q_f(t)$$

Given the differential equation, it is possible to prove (via a Lyapunov function) that the bandwidths satisfy the following constrained optimization problem

Problem

$$\begin{split} & \text{maximize} & \sum_{f \in \mathcal{F}} \frac{\sqrt{2}}{d_f} \arctan\left(\frac{d_f \rho_f}{\sqrt{2}}\right) \\ & \text{subject to} & \forall l \in \mathcal{L}, \sum_{\substack{f \in \mathcal{F} \text{ going through } l}} \rho_f \leq B_l \end{split}$$

- The previous model is in fact for TCP Reno with RED
- With TCP Vegas and Droptail, Low et al. obtain instead

Problem

$$\begin{array}{ll} \text{maximize} & \sum_{f \in \mathcal{F}} d_f \log(\rho_f) \\ \\ \text{subject to} & \forall l \in \mathcal{L}, \sum_{f \in \mathcal{F} \text{ going through } l} \rho_f \leq B_l \end{array}$$

Bottom-up simulation models?

- These models come from the network protocol community
- They're elegant and meant to enlighten us about the fundamental properties of network protocols
- Nowhere in the articles by Low et al. is it suggested to use these models for simulation
- Some parameters may be hard to instantiate for simulation, since users see/care only about the macroscopic behavior
- And in fact, there is a problem with the d_f parameter: equilibrium RTT...

Bottom-up simulation models?

- These models come from the network protocol community
- They're elegant and meant to enlighten us about the fundamental properties of network protocols
- Nowhere in the articles by Low et al. is it suggested to use these models for simulation
- Some parameters may be hard to instantiate for simulation, since users see/care only about the macroscopic behavior
- And in fact, there is a problem with the d_f parameter: equilibrium RTT...

Equilibrium RTT???

- The assumption by Low et al. is that d_f is constant
- What this means is that d_f is constant for a given traffic, i.e., a set of flows and their transfer rates
- But if the transfer rates are computed according to the model (for simulation purposes), then we have a circular dependency
 - \blacksquare d_f tells us what the traffic looks like
 - The traffic defines the value of d_f
- \blacksquare Question: how can we pick d_f for simulation purposes?
- lacksquare Our best guess: $d_f = \sum_{l \in \mathcal{L} \text{ traversed by } l} L_l$
 - \blacksquare L_l : latency of link l
 - We can guess, we're not network protocol people ⊕

Equilibrium RTT???

- The assumption by Low et al. is that d_f is constant
- What this means is that d_f is constant for a given traffic, i.e., a set of flows and their transfer rates
- But if the transfer rates are computed according to the model (for simulation purposes), then we have a circular dependency
 - \blacksquare d_f tells us what the traffic looks like
 - The traffic defines the value of d_f
- **Question:** how can we pick d_f for simulation purposes?
- lacksquare Our best guess: $d_f = \sum_{l \in \mathcal{L} \text{ traversed by } l} L_l$
 - \blacksquare L_l : latency of link l
 - We can guess, we're not network protocol people ©

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

A top-down bandwidth sharing model

- With top-down modeling one hypothesizes a bandwidth sharing model that seems reasonable, and then one attempts to prove it is valid
- Such a model is Max-Min fairness:

Problem

$$\begin{array}{ll} \text{maximize} & \min_{f \in \mathcal{F}} \rho_f \\ \text{subject to} & \forall l \in \mathcal{L}, \sum_{f \ \in \ \mathcal{F} \ \text{going through} \ l} \rho_f \leq B_l \end{array}$$

Makes the least happy flow as happy as possible



A top-down bandwidth sharing model

- With top-down modeling one hypothesizes a bandwidth sharing model that seems reasonable, and then one attempts to prove it is valid
- Such a model is Max-Min fairness:

Problem

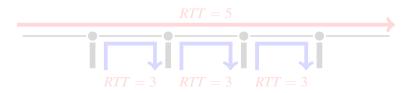
$$\begin{array}{ll} \text{maximize} & \min_{f \in \mathcal{F}} \rho_f \\ \text{subject to} & \forall l \in \mathcal{L}, \sum_{f \ \in \ \mathcal{F} \ \text{going through} \ l} \rho_f \leq B_l \end{array}$$

Makes the least happy flow as happy as possible



Max-min fairness

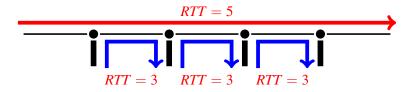
- Max-min fairness is convenient and makes sense
 - what a network should do to be "as fair as possible"
 - Easily computed via a recursive simple algorithm
- Problem: TCP is RTT-unfair



■ On each shared *bottleneck* link: 3/7, 5/7

Max-min fairness

- Max-min fairness is convenient and makes sense
 - what a network should do to be "as fair as possible"
 - Easily computed via a recursive simple algorithm
- Problem: TCP is *RTT-unfair*



■ On each shared bottleneck link: 3/7, 5/7

RTT-unfair Max-min fairness

■ Weighted Max-Min fairness to account for RTT-unfairness

RTT-unfair, windowed Max-min fairness

■ TCP uses a congestion window, *W*, that bounds the number of in-flight bytes

Problem

subject to:

$$\begin{array}{ll} \mathsf{maximize} & \min_{f \in \mathcal{F}} \ L_f \ \rho_f \\ \\ \forall f \in \mathcal{F}, \ L_f = \sum_{l \ \mathsf{traversed} \ \mathsf{by} \ f} \ L_l \\ \\ \forall l \in \mathcal{L}, \ \sum_{f \ \mathsf{going} \ \mathsf{through} \ l} \rho_f \leq B_l \\ \\ \forall f \in \mathcal{F}, \ \rho_f \leq W/(2L_f) \end{array}$$

Validity of the top-down model

- It is known that TCP does not implement (RTT-unfair) Max-min fairness
- But authors have argued it's a reasonable approximation in some cases
 - [Chiu, 1999], [Casanova and Marchal, 2002]
- The big question: is it good enough to be use as a simulation model?
- In [Fujiwara and Casanova, 2007] some good results are presented for a set of topology/flow scenarios
 - But the validity limits are not explored much
 - They do report that for transfers < 100 MiB results are poor (more on this in the next few slides)
- Velho and Legrand went further...

Parameterizing the model

- In [Velho and Legrand, Simutools'09], the authors have evolved this basic model by adding several parameters
- Their approach:
 - Generate a bunch of test topologies
 - Dumbbell topologies with 4 endpoints and 2 routers (>-<)
 - Random with 50 or 200 end-points generated with BRITE
 - Random latencies or latencies computed by BRITE
 - Random bandwidths in various ranges
 - Generate a bunch of flows between random end-points
 - Compare flow-level results to packet-level results
 - Gain insight into what tuning parameters should be added
 - Estimate parameter values of maximum likelihood
- Let's review their findings...

High congestion and RTT-unfairness

- In highly congested scenarios, the RTT is impacted by bandwidth
- New parameter: γ

$$orall f \in \mathcal{F}, \ L_f = \sum_{l \ ext{traversed by } f} L_l + rac{\gamma}{B_l}$$

■ Velho and Legrand find that good results are achieved by $\gamma \approx$ 8775 or $\gamma \approx$ 20537 depending on the TCP version

Protocol overhead

- When using a network protocol, not every transferred byte is a byte of useful data
- Every protocol has an overhead, which decreases the achievable data transfer rate on a link of given bandwidth
- New parameter: β

$$\forall l \in \mathcal{L}, \sum_{f \text{ going through } l} \rho_f \leq \beta \times B_l$$

■ Velho and Legrand find that good results are achieved by $\gamma \approx 0.92$ or $\gamma \approx 0.97$ depending on the TCP version (corresponds to the fraction of useful payload)

Simulating finite flows

- Regardless of the steady-state bandwidth sharing model, we want to simulate finite flows
- $lue{}$ We must compute for each flow its execution time t_f
 - lacksquare to advance the simulated time to $t' = \min(t_{new}, \min_f t_f)$
- A simple model:

$$t_f = L_f + S_f/\rho_f$$

- Problem: TCP's slow start behavior!
 - Due to additive increase, the steady-state bandwidth is not reached immediately, even when there is no congestion

Simulating slow-start

- Because of slow-start, a simulation could only be realistic for large transfer sizes (i.e., ≈ 10 MiB)
 - See the results by Fujiwara and Casanova
- In [Velho et al., 2009] the authors propose a simple modification with a new parameter α :

$$t_f = \frac{\alpha}{\alpha} L_f + S_f / \rho_f$$

- The author find that $\alpha \approx 10.40$ or $\alpha \approx 13.01$ leads to good results depending on the TCP version
- \blacksquare This makes the simulations reasonably accurate down to $\sim 100~{\rm KiB}$ data sizes
- For smaller transfers: forget flow-level models and go back to packet-level since each flow is only a few packets!

The full model in [Velho and Legrand, Simutools'09]

Problem

subject to:

$$egin{aligned} orall f \in \mathcal{F}, \ L_f &= \sum_{l ext{ traversed by } f} L_l \ + \ rac{\gamma}{B_l} \end{aligned}$$
 $orall l \in \mathcal{L}, \ \sum_{f ext{ going through } l}
ho_f \ \leq \ eta imes B_l$ $orall f \in \mathcal{F}, \
ho_f \ \leq \ W \ / \ (2L_f)$

$$t_f = \alpha L_f + S_f / \rho_f$$

The results in [Velho and Legrand, Simutools'09]

- Drastic improvements are demonstrated over the basic Max-Min
 - The addition of each parameter helps in its own way
- Low discrepancy with packet-level simulation shown on hundreds of test cases
- So overall, a very convincing paper with convincing results
- And yet, occasionally, some flow in some test case will experience a discrepancy by up to a factor 20!
- Question: To which extent is the (modified) Max-Min fairness model valid?

The results in [Velho and Legrand, Simutools'09]

- Drastic improvements are demonstrated over the basic Max-Min
 - The addition of each parameter helps in its own way
- Low discrepancy with packet-level simulation shown on hundreds of test cases
- So overall, a very convincing paper with convincing results
- And yet, occasionally, some flow in some test case will experience a discrepancy by up to a factor 20!
- Question: To which extent is the (modified) Max-Min fairness model valid?

Outline

- 1 Current Simulation Practices
- 2 State of the art
- 3 SIMGRID overview
- 4 SIMGRID and scalability
- 5 Network Simulation
 - Steady-state flow-level bandwidth sharing models
 - Bottom-up modeling
 - Top-down modeling
 - On the (top-down) invalidation path

The critical method

- By contrast with bottom-up modeling, top-down modeling is much more of an empirical science
 - (Re)formulate a model as an hypothesis
 - Try to invalidate it via experiments
 - Repeat
- This is the critical method [Karl Popper, philosopher]
 - Don't keep showing "good" cases in which the model works (although it's the typically publication strategy!)
 - Instead keep looking for "bad" cases to evolve the models
 - These bad cases are called *crucial experiments* (i.e., invalidation experiments)
- Invalidation should be at the core of the modeling activity

The critical method

- By contrast with bottom-up modeling, top-down modeling is much more of an empirical science
 - (Re)formulate a model as an hypothesis
 - Try to invalidate it via experiments
 - Repeat
- This is the critical method [Karl Popper, philosopher]
 - Don't keep showing "good" cases in which the model works (although it's the typically publication strategy!)
 - Instead keep looking for "bad" cases to evolve the models
 - These bad cases are called *crucial experiments* (i.e., invalidation experiments)
- Invalidation should be at the core of the modeling activity

Invalidation

- Let us follow the critical method to invalidate the flow-level model in [Velho and Legrand, 2009]
- Essentially, we try to break the model
- So that we can either:
 - Improve it, or
 - Quantify validity limits
- All results are for TCP Reno + RED
 - Some results are discussed with TCP Vegas + Droptail in our recent paper in IEEE TOMACS

Experimental procedure in [Velho and Legrand]

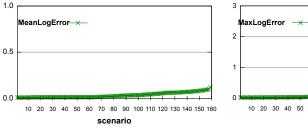
- The results in [Velho and Legrand] are for 4 sets of 10 topologies generated with BRITE, using the Waxman model
 - Number of nodes: 50 or 200
 - Link bandwidths: uniformly distributed in [100, 128] MiB/s or [10, 128] MiB/s
 - Link latencies: computed by BRITE based on Euclidian distance
- For each of the topologies:
 - 100 flows are generated between random end-points
 - Each flow transfers 100 MiB of data
 - TCP congestion window is set to 60 KiB
- In total: 160 experimental scenarios (4×40)

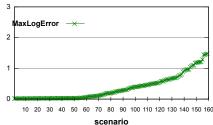
Accuracy metric

$$\begin{split} \mathsf{MeanLogErr} &= \frac{1}{|\mathcal{F}|} \sum_{f \in \mathcal{F}} \left| \log \left(\rho_f^{\mathit{flow}} \right) - \log \left(\rho_f^{\mathit{packet}} \right) \right| \;, \mathsf{and} \\ \mathsf{MaxLogErr} &= \max_{f \in \mathcal{F}} \left| \log \left(\rho_f^{\mathit{flow}} \right) - \log \left(\rho_f^{\mathit{packet}} \right) \right| \;. \end{split}$$

- The smaller the error, the better
- The logarithms are used for symmetry (doesn't matter whether the "reference" is the smaller or the larger value)
- Computed over the time window from t = 0 up to the completion of the first flow

Initial results



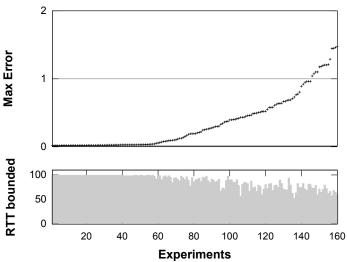


- A few scenarios show high error
- Many show low error: they're not "crucial experiments"

Making scenarios more crucial

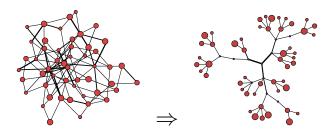
- Inspecting the results, we find that many of the low-error scenarios have mostly RTT-bound flows
- They are governed by the $\rho_f \leq W/(2L_f)$ constraint
 - The "real" bandwidth sharing model is bypassed
- They are easy cases for our flow-level model, hence the low error
- We can see this trend easily if we plot % of RTT-bounded flow vs. MaxError...

More RTT-bounded flow, less crucial

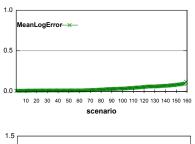


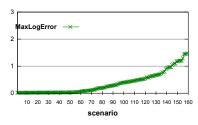
Making experiments more critical

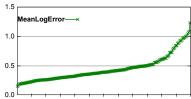
- We sample bandwidths in [10, 12.8] MiB/s
 - Likely rare in practice, but we use the critical method
- We no longer use BRITE, but instead use Tiers

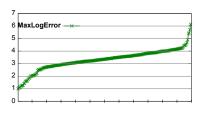


New crucial workload





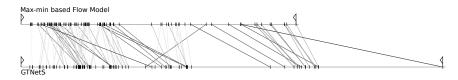




Finding sources of error

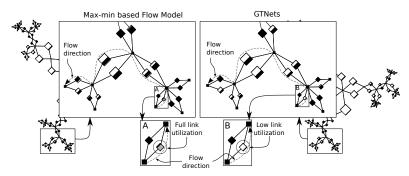
- Now we have a lot of bad results, which is great because now it's easier to try to understand what's wrong with the model
- We pick a particularly bad workload and visualize flow completion time in flow-level simulation and packet-level simulation...

Visualization of completion times



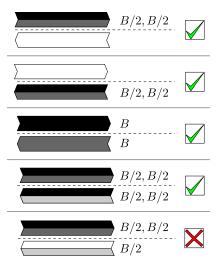
- Many almost vertical lines (light)
- Several really non-vertical lines (dark)
- Flow-level simulation almost always underestimates completion times
 - We have no idea how to explain the one odd flow that is faster with packet-level simulation
- We now use another visualization to inspect the "bad" flows

Visualization of link usage



■ There is no hope for our flow-level model to capture the underutilization of that link...

Link underutilization explained



- Underutilization is due to reverse traffic
- One explanation: Network compression [Zhang 1991]
 - Ack packets are queued with the data packets of reverse flows
- Another: Ack-clocking [Heusse, 2011]
 - Data and ack segments alternatively fill only one link buffer
- We modify our problem...

Accounting for reverse traffic in our model

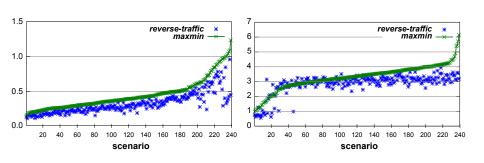
Problem

subject to:

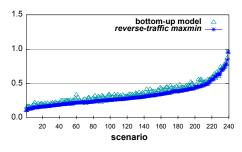
$$orall f \in \mathcal{F}, \; L_f = \sum_{l \; ext{traversed by} \; f} \; L_l \; + \; rac{\gamma}{B_l}$$

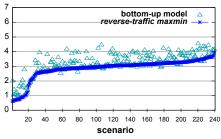
$$egin{aligned} orall l \in \mathcal{L}, & \sum_{f ext{ going through } l}
ho_f + \epsilon \left(\sum_{f ext{'s ack going through } l}
ho_f
ight) \leq \ eta imes B_l \ orall f \in \mathcal{F}, &
ho_f \ \leq \ W \ / \ (2L_f) \ & t_f \ = \ lpha \ L_f \ + \ S_f \ / \
ho_f \end{aligned}$$

Improved results



Comparison to bottom-up model





Could bottom-up be improved?

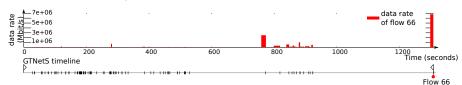
- The model by Low et al. does not account for reverse traffic
- Since the work by Low et al., new advances have been made
 - Differential Algebraic Equation (DAE) model in [Tang et al., 2008, 20010]
 - Ack-clocking dynamics in [Jacobsson et al., 2009]
- These models are not designed for use in simulation
 - They require solving complex differential equations
 - Models not yet validated for more than three flows and three nodes
- Conclusion: At this point, our top-down modified Max-Min model looks pretty good

Could bottom-up be improved?

- The model by Low et al. does not account for reverse traffic
- Since the work by Low et al., new advances have been made
 - Differential Algebraic Equation (DAE) model in [Tang et al., 2008, 20010]
 - Ack-clocking dynamics in [Jacobsson et al., 2009]
- These models are not designed for use in simulation
 - They require solving complex differential equations
 - Models not yet validated for more than three flows and three nodes
- Conclusion: At this point, our top-down modified Max-Min model looks pretty good

Some things, we just can't do...

 Some flow just experience data rates that cannot be modeled by a flow-level model



- This flows stalls for 380 seconds!
 - Results reproduced with NS-3 and GTNetS
- This is for a very (unlikely) high-contention scenario

What we found out

- Flow-level models used in state-of-the-art grid/cloud simulators can be vastly improved upon
- Results are close to packet-level results unless
 - Transferred data sizes are small (a few KiB)
 - Congestion is really high
- Using the critical method was key to obtained improvements in the model
- We feel that we have reached the limits of what we can do